

Exercise - Integrate a Notebook within Azure Synapse Pipelines

In this unit, you create an Azure Synapse Spark notebook to analyze and transform data loaded by a mapping data flow and store the data in a data lake. You create a parameter cell that accepts a string parameter that defines the folder name for the data the notebook writes to the data lake. You then add this notebook to a Synapse pipeline and pass the unique pipeline run ID to the notebook parameter so that you can later correlate the pipeline run with the data saved by the notebook activity. Finally, you use the Monitor hub in Synapse Studio to monitor the pipeline run, obtain the run ID, then locate the corresponding files stored in the data lake.

About Apache Spark and notebooks

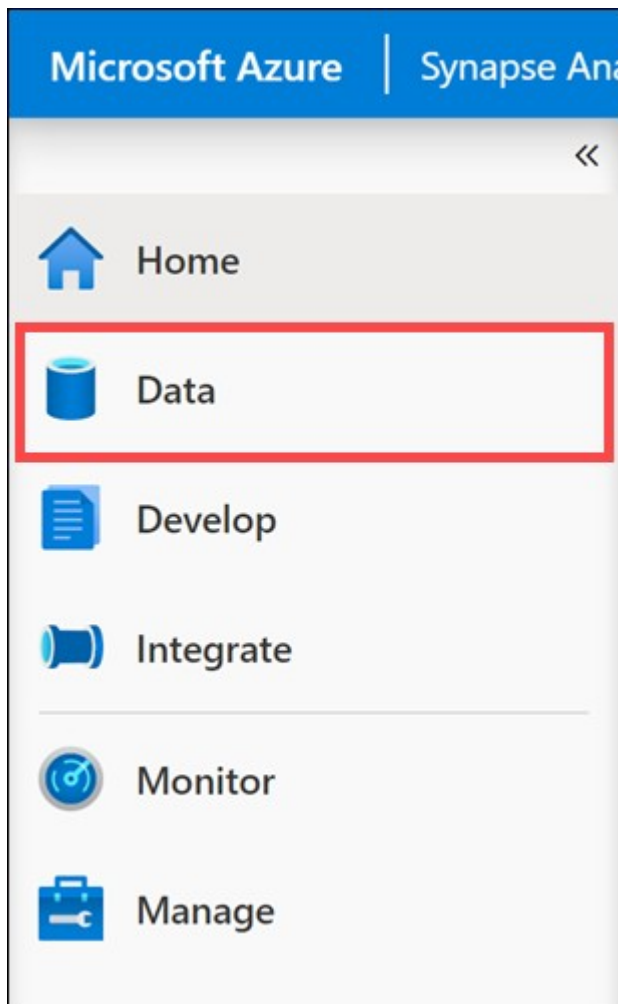
Apache Spark is a parallel processing framework that supports in-memory processing to boost the performance of big-data analytic applications. Apache Spark in Azure Synapse Analytics is one of Microsoft's implementations of Apache Spark in the cloud.

An Apache Spark notebook in Synapse Studio is a web interface for you to create files that contain live code, visualizations, and narrative text. Notebooks are a good place to validate ideas and use quick experiments to get insights from your data. Notebooks are also widely used in data preparation, data visualization, machine learning, and other Big Data scenarios.

Create a Synapse Spark notebook

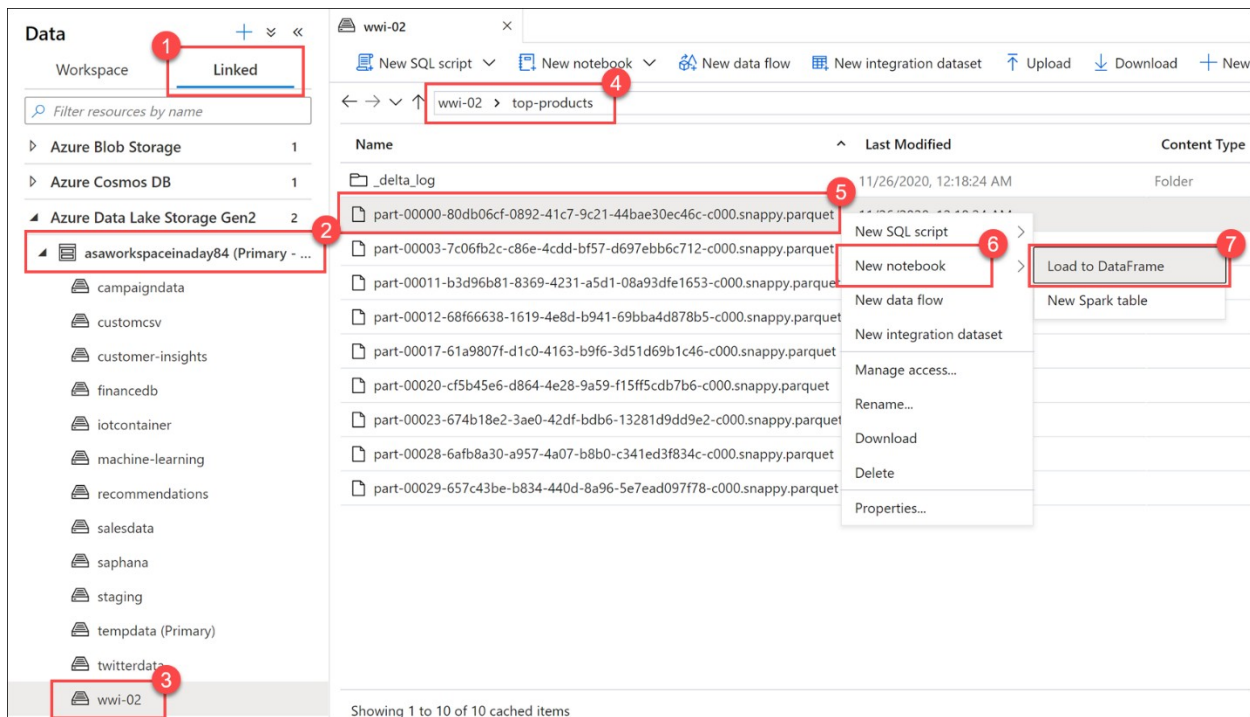
Suppose you created a Mapping Data flow in Synapse Analytics to process, join, and import user profile data. Now you want to find the top 5 products for each user, based on which ones are both preferred and top, and have the most purchases in the past 12 months. Then, you want to calculate the top 5 products overall.

In this step, you create a Synapse Spark notebook to make these calculations. 1. Open [Synapse Analytics Studio](#), and then navigate to the Data hub.



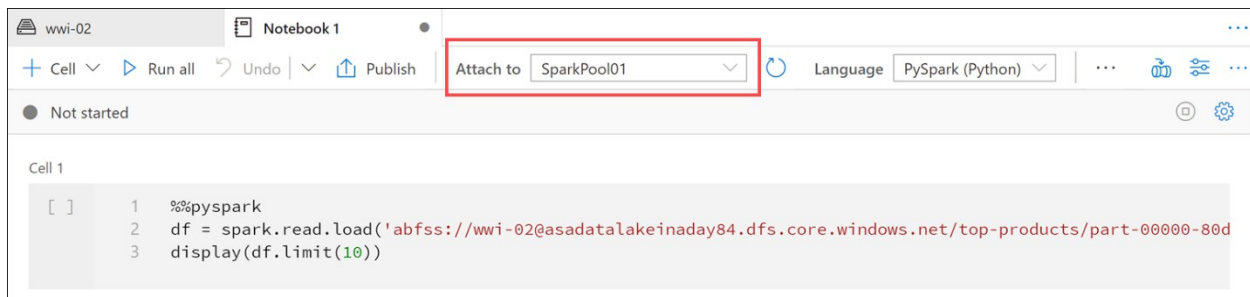
Data hub

2. Select the **Linked** tab (1) and expand the **primary data lake storage account** (2) underneath the **Azure Data Lake Storage Gen2**. Select the **wwi-02** container (3) and open the **top-products** folder (4). Right-click on any Parquet file (5), select the **New notebook** menu item (6), then select **Load to DataFrame** (7). If you don't see the folder, select Refresh above.



The Parquet file and new notebook option are highlighted

3. Make sure the notebook is attached to your Spark pool.



The attach to Spark pool menu item is highlighted

4. Replace the Parquet file name with **.parquet (1)* to select all Parquet files in the *top-products* folder. For example, the path should be similar to:

abfss://wwi-02@YOUR_DATALAKE_NAME.dfs.core.windows.net/top-products/*.parquet.



The filename is highlighted.

5. Select **Run all** on the notebook toolbar to execute the notebook.

Cell 1

```
1 %%pyspark
2 df = spark.read.load('abfss://wwi-02@asadatalakeinaday84.dfs.core.windows.net/top-products/*.parquet')
3 display(df.limit(10))
```

Command executed in 2mins 35s 588ms by joel on 11-26-2020 00:53:24.571 -05:00

> **Job execution** Succeeded **Spark** 2 executors 8 cores [View in monitoring](#) [Open Spark UI](#)

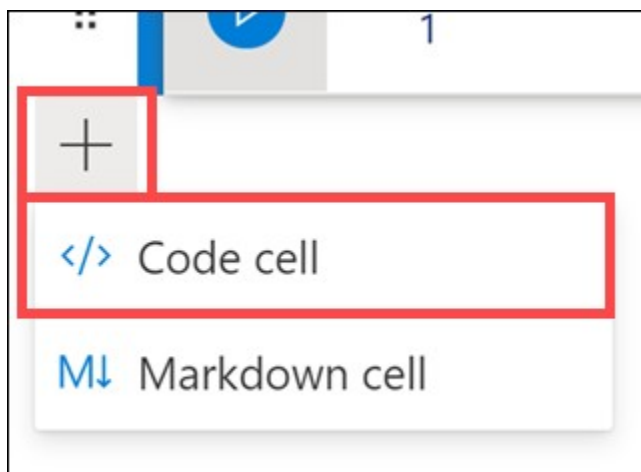
View **Table** Chart

visitorId	productId	itemsPurchased...	preferredProduc...	userId	isTopProduct	isPreferredProd...
	2717		2717	148	false	true
	4002		4002	148	false	true
	1716		1716	148	false	true
	4520		4520	148	false	true
	951		951	148	false	true
	1817		1817	148	false	true
	2634		2634	463	false	true
	2795		2795	463	false	true

The cell results are displayed.

Note: The first time you run a notebook in a Spark pool, Synapse creates a new session. This can take approximately 3-5 minutes. **Note:** To run just the cell, either hover over the cell and select the *Run cell* icon to the left of the cell, or select the cell then type **Ctrl+Enter** on your keyboard.

6. Create a new cell underneath by selecting the + button and selecting the **Code cell** item. The + button is located beneath the notebook cell on the left. Alternatively, you can also expand the + **Cell** menu in the Notebook toolbar and select the **Code cell** item.



The Add Code menu option is highlighted.

7. Enter and execute the following in the new cell to populate a new dataframe called *top_purchases*, create a new temporary view named *top_purchases* , and show the first 100 rows:

```
topPurchases = df.select(
    "UserId", "ProductId",
    "ItemsPurchasedLast12Months", "IsTopProduct",
    "IsPreferredProduct")

# Populate a temporary view so we can query from SQL
topPurchases.createOrReplaceTempView("top_purchases")

topPurchases.show(100)
```

The output should look similar to the following:

```
| 833| 1087| null| false| true|
```

8. Execute the following in a new cell to create a new temporary view by using SQL:

```
%%sql
```

```
CREATE OR REPLACE TEMPORARY VIEW top_5_products
```

AS

```
select UserId, ProductId, ItemsPurchasedLast12Months
from (select *,
            row_number() over (partition by UserId order by ItemsPurchasedLa
st12Months desc) as seqnum
      from top_purchases
     ) a
where seqnum <= 5 and IsTopProduct == true and IsPreferredProduct = true
order by a.UserId
```

Note that there is no output for the above query. The query uses the `top_purchases` temporary view as a source and applies a `row_number()` over method to apply a row number for the records for each user where `ItemsPurchasedLast12Months` is greatest. The `where` clause filters the results so we only retrieve up to five products where both `IsTopProduct` and `IsPreferredProduct` are set to true. This gives us the top five most purchased products for each user where those products are also identified as their favorite products, according to their user profile stored in Azure Cosmos DB. 9. Execute the following in a new cell to create and display a new DataFrame that stores the results of the `top_5_products` temporary view you created in the previous cell:

```
top5Products = sqlContext.table("top_5_products")

top5Products.show(100)
```

You should see an output similar to the following, which displays the top five preferred products per user:

1
2
3

Cell 5

```
[26] 1 top5Products = sqlContext.table("top_5_products")
      2
      3 top5Products.show(100)
```

UserId	ProductId	ItemsPurchasedLast12Months
80000	2069	93
80000	2069	93
80000	2069	93
80000	2069	93
80000	2069	93
80001	1812	93
80001	1812	93
80001	1812	93
80001	1812	93
80001	1812	93
80002	1256	90
80002	1256	90
80002	4987	88
80002	3190	92
80002	3190	92
80003	295	91
80003	638	97

The top five preferred products are displayed per user.

10. Calculate the top five products overall, based on those that are both preferred by customers and purchased the most. To do this, execute the following in a new cell:

```
1
2
3
4
5
6
7
top5ProductsOverall = (top5Products.select("ProductId", "ItemsPurchasedLast12Months")
    .groupBy("ProductId"))
```

```
.agg( sum("ItemsPurchasedLast12Months").alias("Total") )
.orderBy( col("Total").desc() )
.limit(5))
```

```
top5ProductsOverall.show()
```

In this cell, we grouped the top five preferred products by product ID, summed up the total items purchased in the last 12 months, sorted that value in descending order, and returned the top five results. Your output should be similar to the following:

```
1
2
3
4
5
6
7
8
9
```

ProductId	Total
2107	4538
4833	4533
347	4523
3459	4233
4246	4155

Create a parameter cell

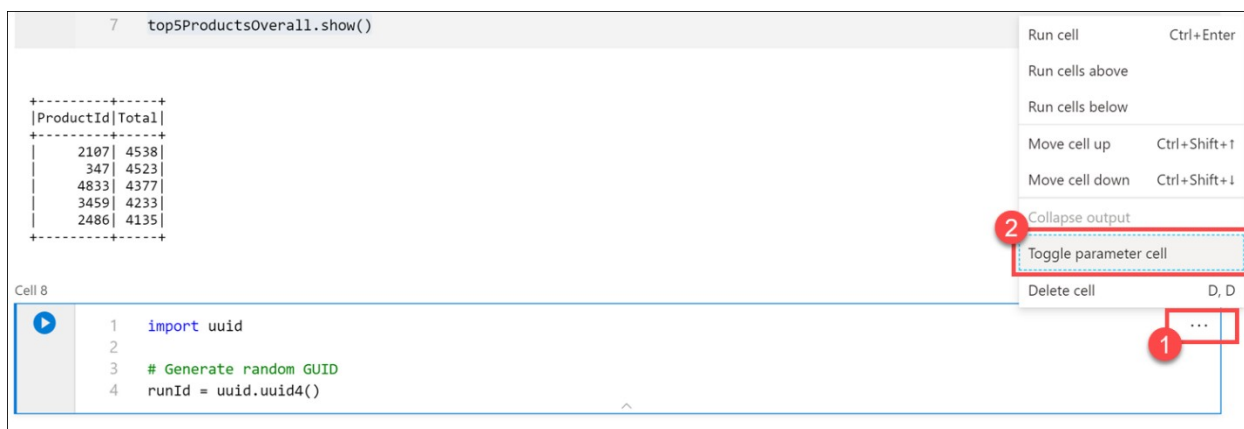
Azure Synapse pipelines look for the parameters cell and treat this cell as defaults for the parameters passed in at execution time. The execution engine will add a new cell beneath the parameters cell with input parameters in order to overwrite the default values. When a parameters cell isn't designated, the injected cell will be inserted at the top of the notebook. 1. We are going to execute this notebook from a pipeline. We want to pass in a parameter that sets a *runId* variable value that will be used to name the Parquet file. Execute the following in a new cell:

2
3
4

```
import uuid

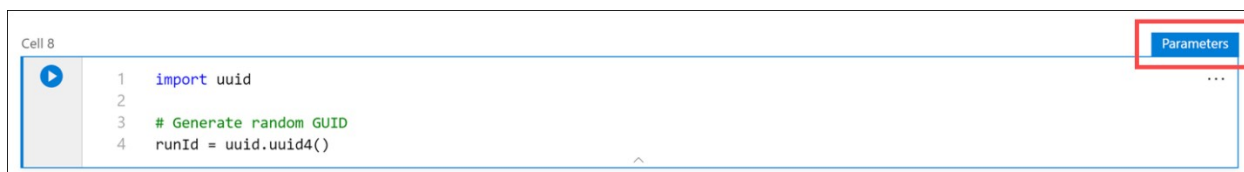
# Generate random GUID
runId = uuid.uuid4()
```

We are using the *uuid* library that comes with Spark to generate a random GUID. We want to override the *runId* variable with a parameter passed in by the pipeline. To do this, we need to toggle this as a parameter cell. 2. Select the actions ellipses (...) on the top-right corner of the cell (1), then select **Toggle parameter cell** (2).



The menu item is highlighted.

After toggling this option, you will see the **Parameters** tag on the cell



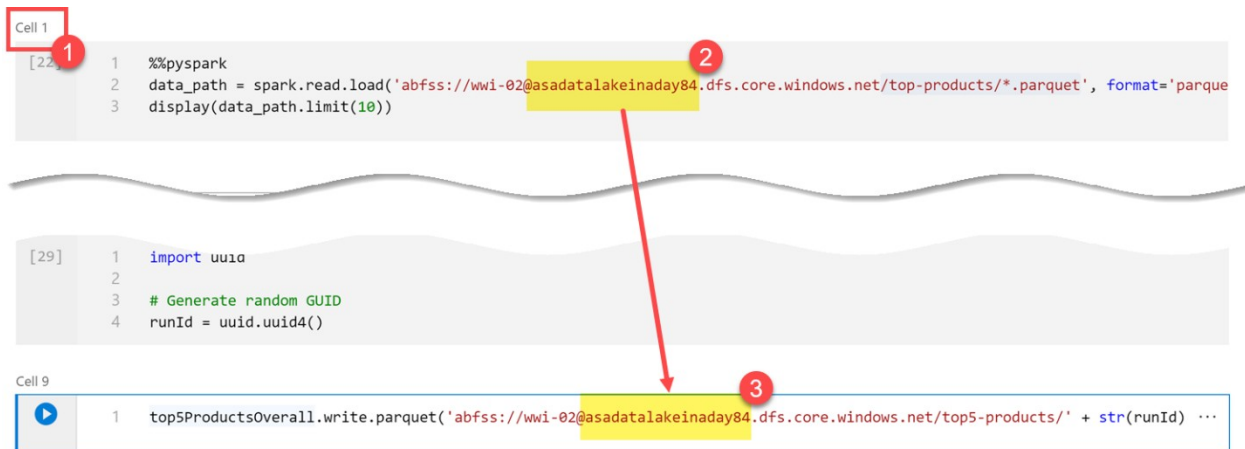
The cell is configured to accept parameters.

3. Paste the following code in a new cell to use the *runId* variable as the Parquet filename in the */top5-products/* path in the primary data lake account. **Replace** YOUR DATALAKE NAME in the path with the name of your primary data lake account. To find this, scroll up to **Cell 1** at the top of the page (1). Copy the data lake storage account from the path (2). Paste this value as a replacement for YOUR DATALAKE NAME in the path (3) inside the new cell, then execute the cell

1
2
3

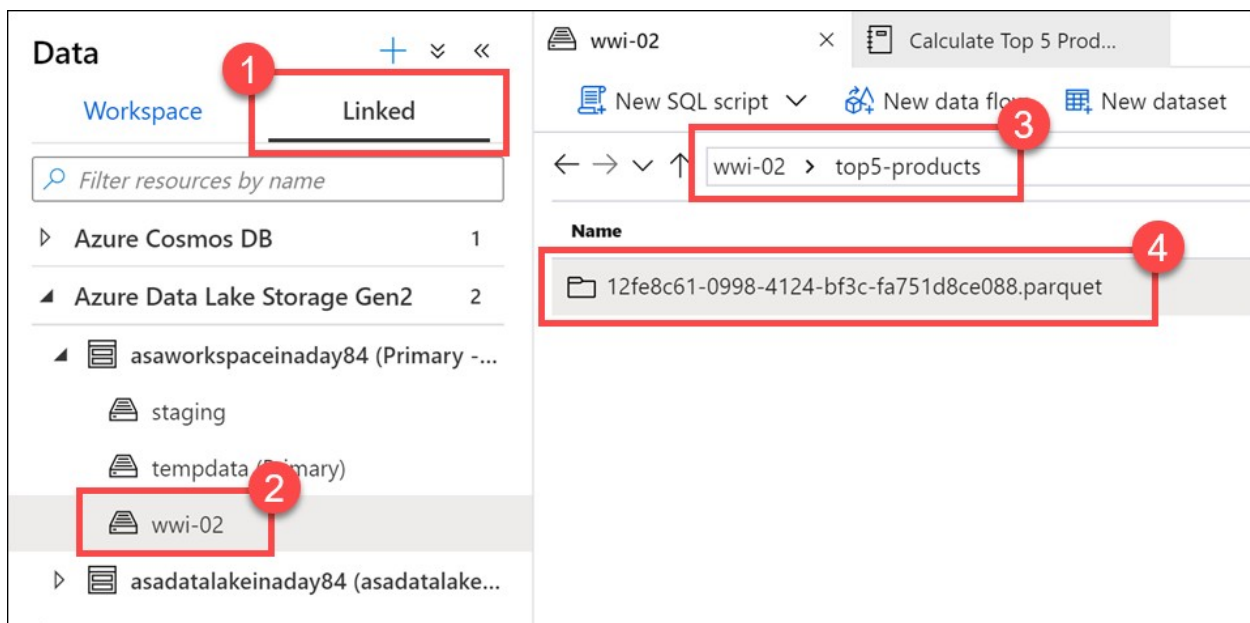
```
%%pyspark
```

```
top5ProductsOverall.write.parquet('abfss://wwi-02@YOURDATA LAKE_NAME.dfs.core.windows.net/top5-products/' + str(runId) + '.parquet')
```



The path is updated with the name of the primary data lake account

4. Verify that the file was written to the data lake. Navigate to the **Data** hub and select the **Linked** tab (1). Expand the primary data lake storage account and select the **wwi-02** container (2). Navigate to the **top5-products** folder (3). You should see a folder for the Parquet file in the directory with a GUID as the file name (4).

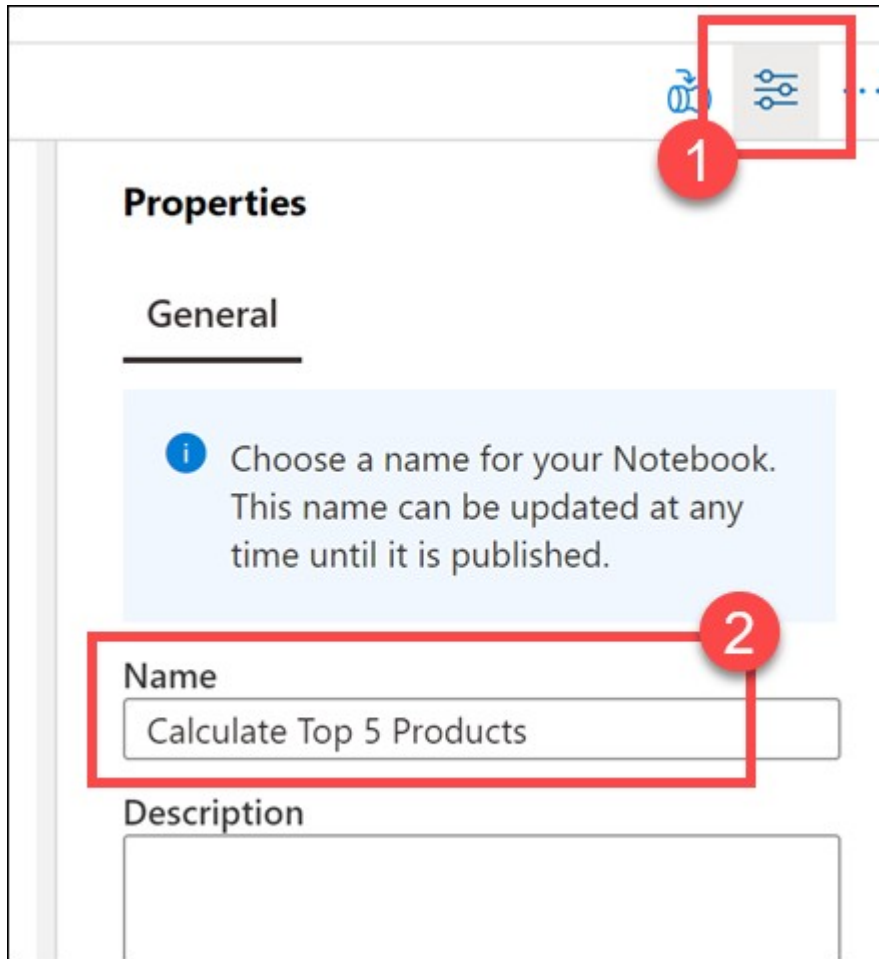


The parquet file is highlighted.

The Parquet write method on the dataframe in the Notebook cell created this directory since it did not previously exist. Add the Notebook to a Synapse pipeline

Referring back to the Mapping Data Flow we mentioned at the beginning of the exercise, suppose you want to execute this notebook after the Data Flow runs as part of your orchestration process. To do this, you add this notebook to a pipeline as a new Notebook activity.

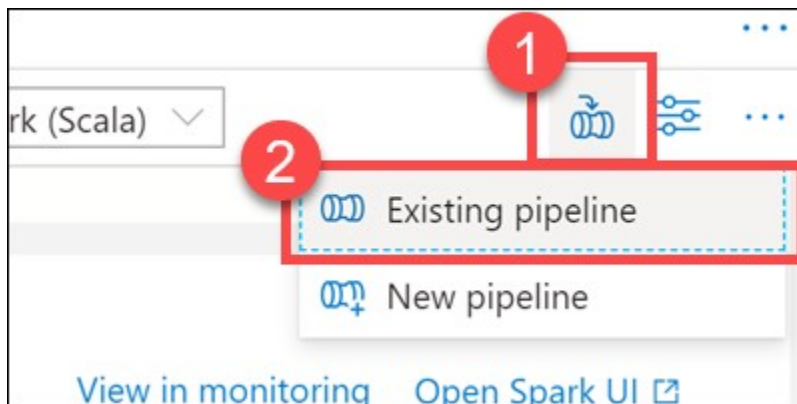
1. Return to the notebook. Select the **Properties** button (1) at the top-right corner of the notebook, then enter *Calculate Top 5 Products* for the Name (2).



The screenshot shows the 'Properties' blade with the 'General' tab active. A red box labeled '1' points to the 'Properties' button in the top-right corner. Another red box labeled '2' points to the 'Name' input field, which contains the text 'Calculate Top 5 Products'.

The properties blade is displayed.

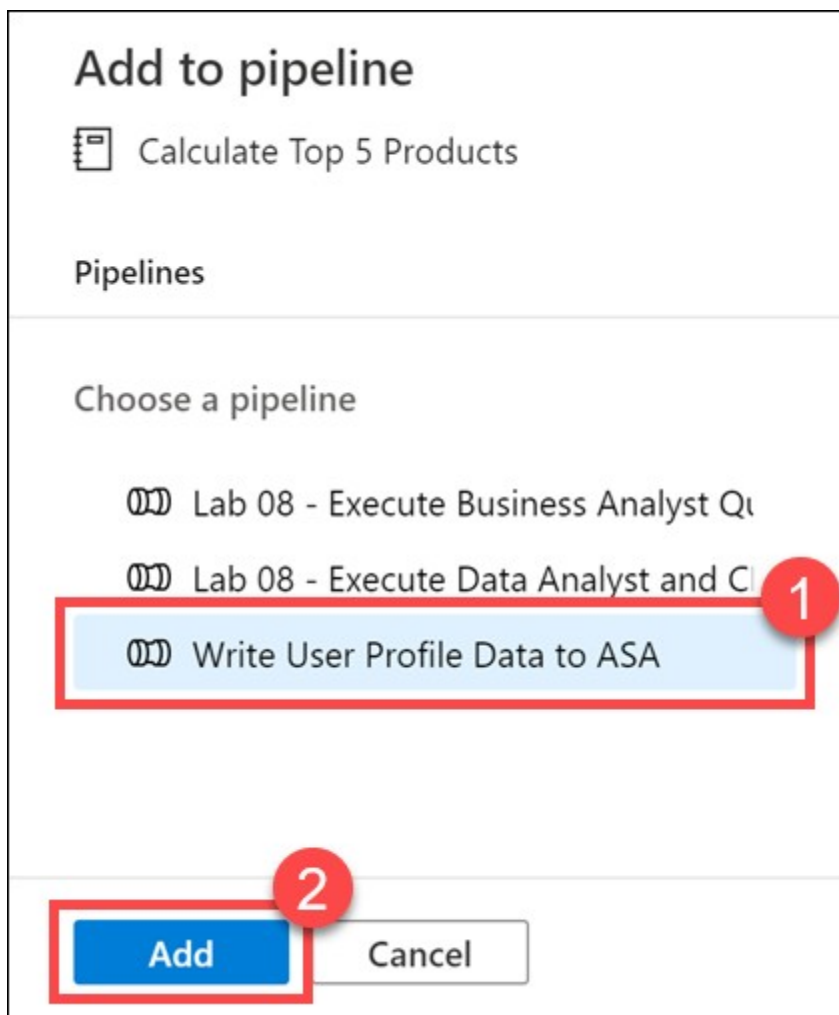
2. Select the **Add to pipeline** button (1) at the top-right corner of the notebook, then select **Existing pipeline** (2)



The add to pipeline button is

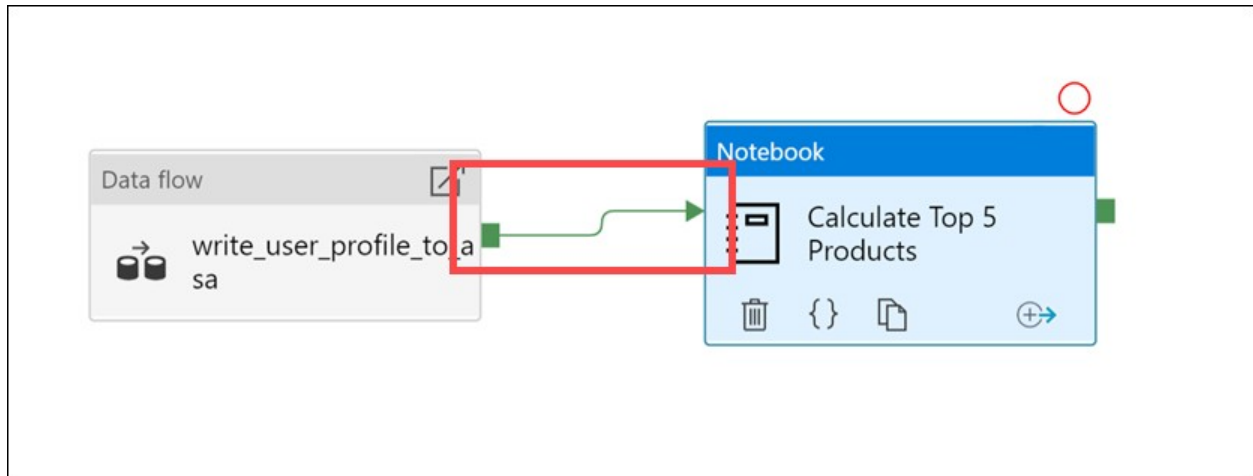
highlighted.

3. Select the Write User Profile Data to ASA pipeline (1), then select Add *2).



The pipeline is selected.

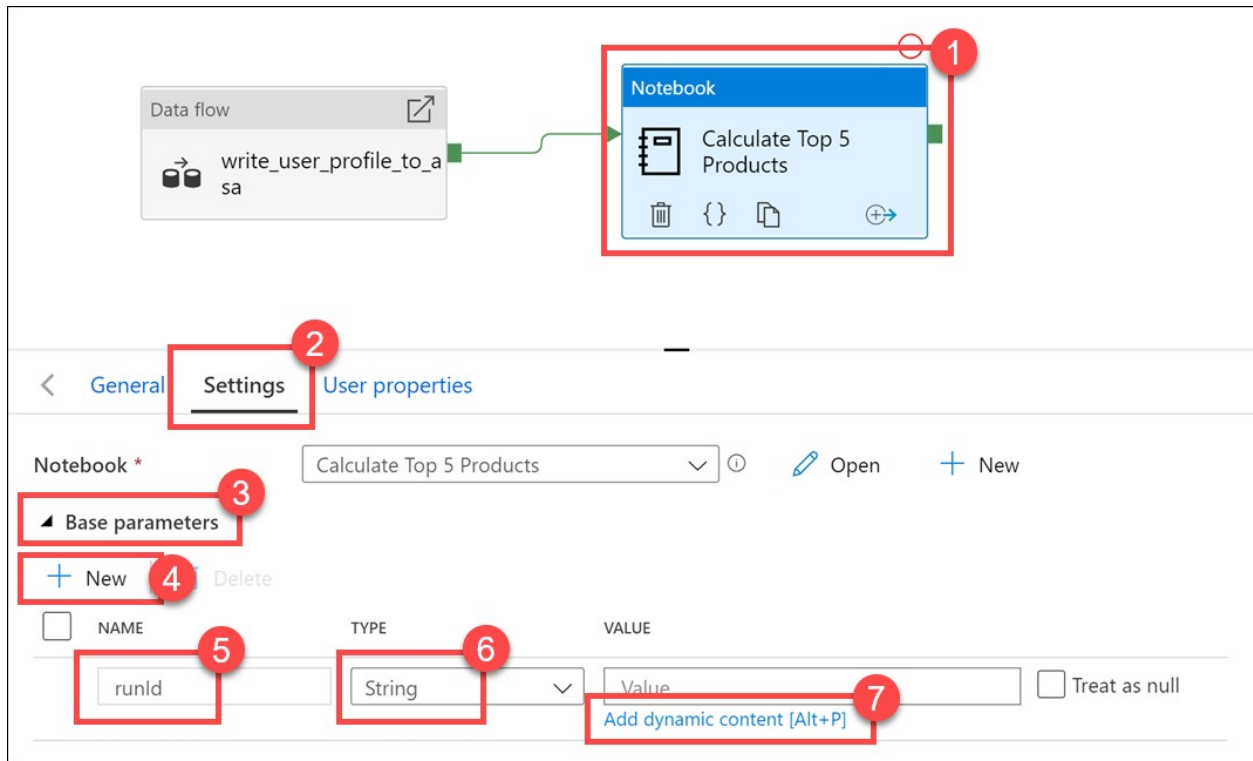
4. Synapse Studio adds the Notebook activity to the pipeline. Rearrange the **Notebook** activity so it sits to the right of the **Data flow** activity. Select the **Data flow** activity and drag a **Success** activity pipeline connection green box to the **Notebook** activity.



The green arrow is highlighted.

The **Success** activity arrow instructs the pipeline to execute the **Notebook** activity after the **Data flow** activity successfully runs.

5. Select the **Notebook** activity (1), select the **Settings** tab (2), expand **Base parameters** (3), and select **+ New** (4). Enter *runId* in the **Name** field (5). Select **String** for the **Type** (6). For the **Value**, select **Add dynamic content** (7)



The settings are displayed.

6. Select Pipeline run ID under System variables (1). This adds `@pipeline().RunId` to the dynamic content box (2). Select Finish (3) to close the dialog.

Add dynamic content

@pipeline().RunId 2

[Clear contents](#)

[Filter...](#)



Use [expressions](#), [functions](#) or refer to [system variables](#).

System variables

Pipeline Name

Name of the pipeline

Pipeline run ID

ID of the specific pipeline run

Pipeline trigger ID

ID of the trigger that invokes the pipeline

Pipeline trigger name

Name of the trigger that invokes the pipeline

Pipeline trigger time

Time when the trigger that invoked the pipeline. The trigger time is the actual fired time, not the sc...

Pipeline trigger type

Type of the trigger that invoked the pipeline (Manual, Scheduler)

Workspace name

Name of the workspace the pipeline run is running within

Functions

✖ [Expand all](#)

▸ Collection Functions

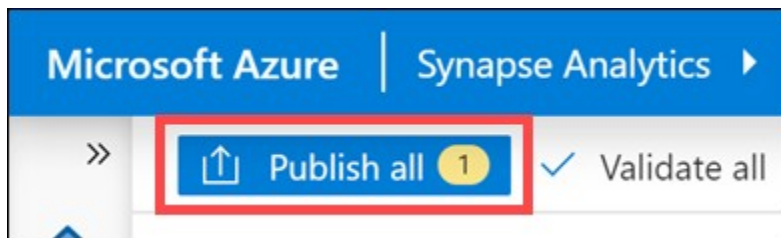
Finish 3

Cancel

The dynamic content form is displayed.

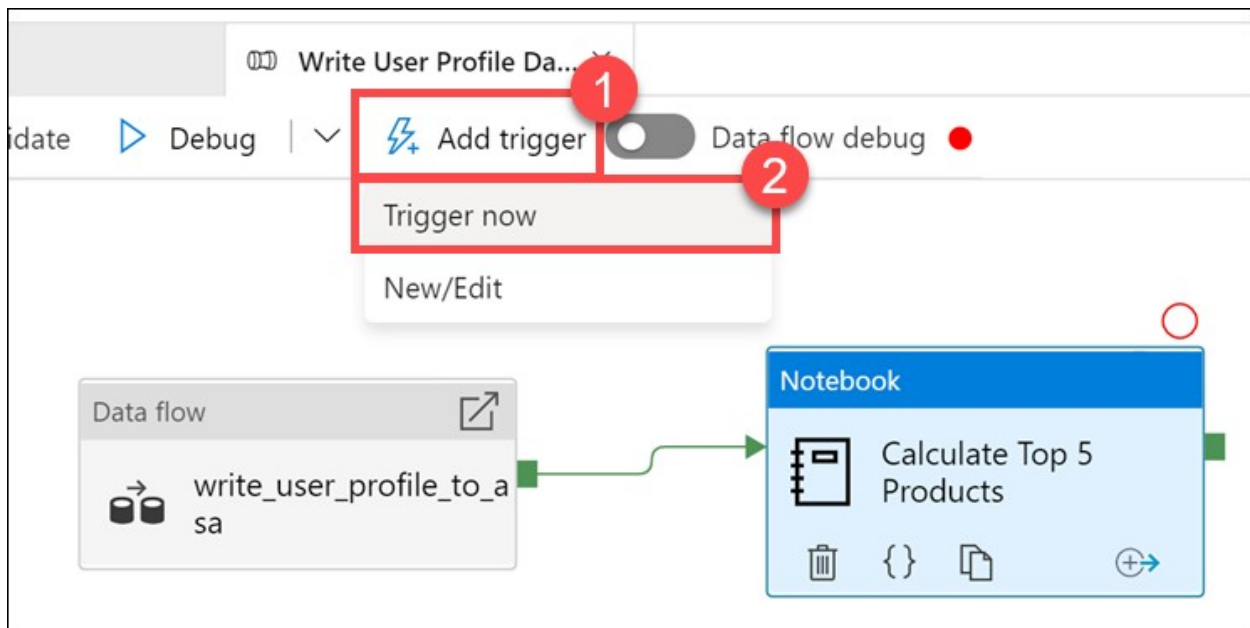
The Pipeline run ID value is a unique GUID assigned to each pipeline run. We will use this value for the name of the Parquet file by passing this value in as the *runId* Notebook parameter. We can then look through the pipeline run history and find the specific Parquet file created for each pipeline run.

7. Select **Publish all** then **Publish** to save your changes



Publish all is highlighted.


8. After publishing is complete, select **Add trigger (1)**, then **Trigger now (2)** to run the updated pipeline.



The trigger menu item is highlighted.

9. Select **OK** to run the trigger.

Pipeline run

 Trigger pipeline now using last published configuration.

Parameters

NAME	TYPE	VALUE
No records found		

OK

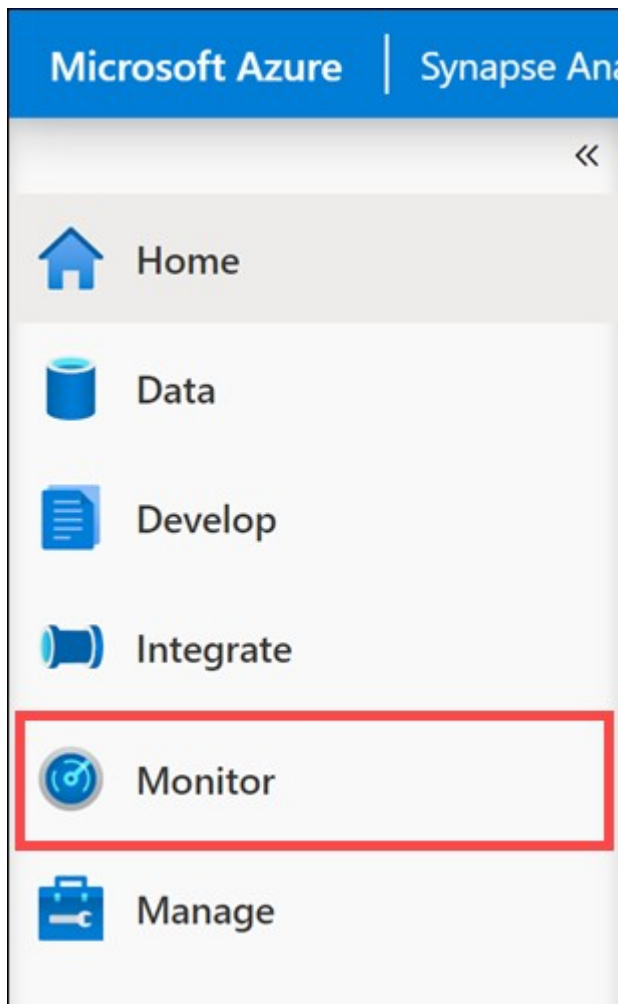
Cancel

The OK button is highlighted.

Monitor the pipeline run

The Monitor hub lets you monitor current and historical activities for SQL, Apache Spark, and Pipelines.

1. Navigate to the **Monitor** hub.



The Monitor hub menu item is selected.

2. Select Pipeline runs (1) and wait for the pipeline run to successfully complete (2). You may need to refresh (3) the view.




The screenshot displays the 'Pipeline runs' section of the Azure Synapse Analytics interface. On the left, the 'Integration' sidebar has 'Pipeline runs' selected, marked with a red circle and the number 1. The main area shows a table of pipeline runs. At the top of this area, the 'Refresh' button is highlighted with a red circle and the number 3. In the table, the 'Status' column for the first two rows shows 'Succeeded', with the first instance highlighted by a red circle and the number 2. The table includes columns for Pipeline name, Run start, Run end, Duration, Triggered by, Status, and Run.


Pipeline name	Run start	Run end	Duration	Triggered by	Status	Run
Write User Profile Data t...	11/26/20, 3:33:48 AM	11/26/20, 3:45:43 AM	00:11:55	Manual trigger	Succeeded	Original
Write User Profile Data t...	11/26/20, 3:16:33 AM	11/26/20, 3:21:26 AM	00:04:53	Manual trigger	Succeeded	Original

The pipeline run succeeded.


3. Select the name of the pipeline to view the pipeline's activity runs.

Pipeline runs


Triggered Debug |  Rerun  Cancel  Refresh

 Search by run ID or name

Eastern Time (US & C... : **Last 24 h**

 Add filter

Showing 1 - 11 items

<input type="checkbox"/> Pipeline name	Run start 
<input type="checkbox"/> Write User Profile Data to ASA	11/26/20, 3:33:48 AM

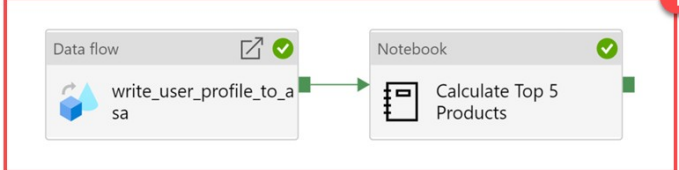
The pipeline name is selected.

4. Notice both the **Data flow** activity, and the new **Notebook** activity (1). Make note of the **Pipeline run ID** value (2). We will compare this to the Parquet file name generated by the notebook. Select the **Calculate Top 5 Products Notebook** name to view its details (3).

Write User Profile Data to ASA

List Gantt

Rerun Rerun from activity Rerun from failed activity Refresh



Activity runs

Pipeline run ID 16bb8447-824a-4147-b7ea-348505ebdc44

All status ▾

Showing 1 - 2 of 2 items

Activity name	Activity type	Run start ↑↓	Duration	Status	Integration runtime
Calculate Top 5 Products	SynapseNotebo	11/26/20, 3:39:57 AM	00:05:45	✓ Succeeded	DefaultIntegrationRuntime (East US 2)
write_user_profile_to_asa	ExecuteDataFlov	11/26/20, 3:33:50 AM	00:06:07	✓ Succeeded	DefaultIntegrationRuntime (East US 2)

The pipeline run details are displayed.

5. Here we see the Notebook run details. You can select the **Playback button (1) to watch a playback of the progress through the **jobs** (2). At the bottom, you can view the **Diagnostics** and **Logs** with different filter options (3). To the right, we can view the run details, such as the duration, Livy ID, Spark pool details, etc. Select the **View details** link on a **job** to view its details (5).**

bf5a1022-c80d-4f98-b60f-7b2ce74606ff
Completed tasks 2562 of 2562 Status Stopped Total duration 5m 45s

Cancel Refresh Spark history server

Attempts 1 of 1
All job IDs View Progress Playback 0ms / 2min 5s 494ms

Stage 0
Job 0
Tasks: 1
Duration: 6s 510ms
Rows: 0
Data read: 0Byte
Data written: 0Byte
[View details](#)

Stage 1
Job 1
Tasks: 1
Duration: 6s 892ms
Rows: 4096
Data read: 1.6MBs
Data written: 0Byte
[View details](#)

Stage 2
Job 2
Tasks: 8
Duration: 2s 469ms
Rows: 162203
Data read: 5.7MBs
Data written: 0Byte
[View details](#)

Stage 3
Job 3
Tasks: 8
Duration: 2s 539ms
Rows: 324406
Data read: 5.7MBs
Data written: 10.1MBs
[View details](#)

Stage 4
Job 4
Tasks: 200
Duration: 3s 90ms
Rows: 162203
Data read: 10.1MBs
Data written: 0Byte
[View details](#)

Stage 5
Job 5
Tasks: 8
Duration: 680ms
Rows: 162203
Data read: 875.3KBs
Data written: 0Byte
[View details](#)

Stage 6
Job 6
Tasks: 8
Duration: 680ms
Rows: 162203
Data read: 875.3KBs
Data written: 0Byte
[View details](#)

Stage 7
Job 7
Tasks: 8
Duration: 680ms
Rows: 162203
Data read: 875.3KBs
Data written: 0Byte
[View details](#)

Diagnostics Logs

- > Failed jobs
- > Data skew
- > Time skew
- > Executor utilization

Details

Summary

Application

Application ID
application_1606380084300_0001

Queued duration
0s

Running duration
5m 45s

Livy ID
4

Submitter
ee20d9e7-6295-4240-ba3f-c3784616c565

Submit time
11/26/20 3:39:58 AM

Executors
2

Spark pool

Name
SparkPool01

The run details are displayed.

6. The Spark application UI opens in a new tab where we can see the stage details. Expand the **DAG Visualization to view the stage details.**

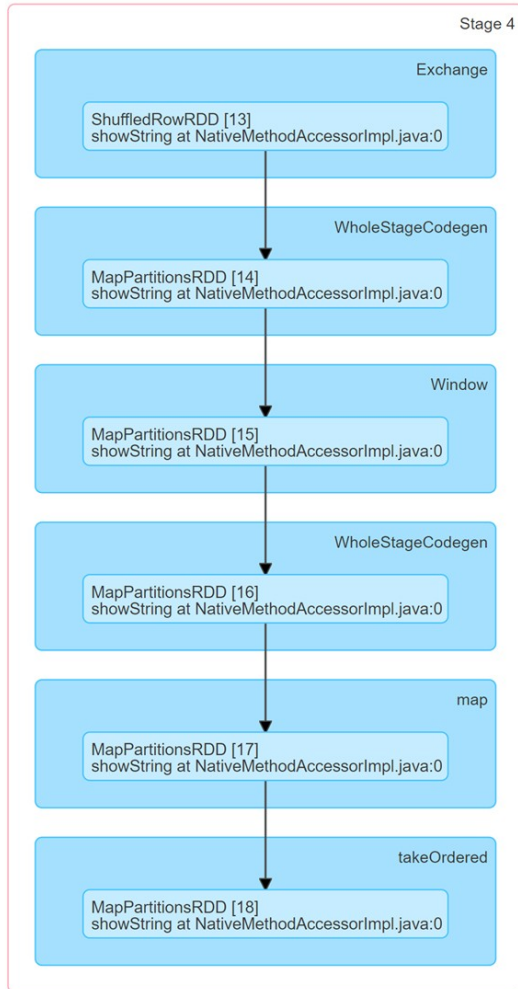
Details for Stage 4 (Attempt 0)

Total Time Across All Tasks: 18 s

Locality Level Summary: Node local: 200

Shuffle Read: 10.1 MB / 1622203

▼ DAG Visualization



► Show Additional Metrics

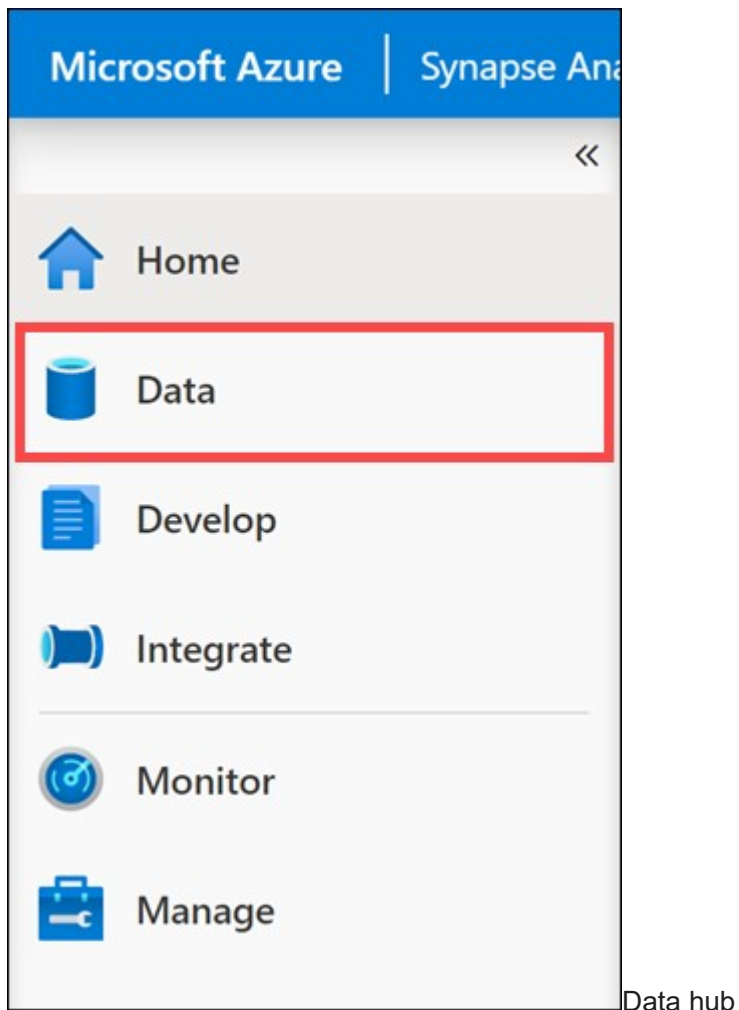
► Event Timeline

Summary Metrics for 200 Completed Tasks

Metric	Min	25th percentile	Median
Duration	20 ms	42 ms	57 ms
GC Time	0 ms	0 ms	0 ms
Shuffle Read Size / Records	43.9 KB / 6471	50.2 KB / 7648	51.5 KB / 8057

The Spark stage details are displayed.

7. Navigate back to the Data hub.



8. Select the **Linked** tab (1), select the **wwi-02** container (2) on the primary data lake storage account, navigate to the **top5-products** folder (3), and verify that a folder exists for the Parquet file whose name matches the **Pipeline run ID**.

The screenshot displays the Azure Data Explorer interface. On the left, the 'Data' pane shows a list of resources under the 'Workspace' tab. The 'Linked' tab is selected (callout 1). The 'wwi-02' folder is highlighted (callout 2). On the right, the 'File Explorer' pane shows the contents of the 'wwi-02' folder. The breadcrumb path 'wwi-02 > top5-products' is highlighted (callout 3). The file '16bb8447-824a-4147-b7ea-348505ebdc44.parquet' is highlighted (callout 4).

The file is highlighted.

As you can see, we have a file whose name matches the Pipeline run ID we noted earlier:

Write User Profile Data to ASA

List Gantt

Rerun Rerun from activity Rerun from failed activity Refresh

Activity runs

Pipeline run ID 16bb8447-824a-4147-b7ea-348505ebdc44

All status ▾

Showing 1 - 2 of 2 items

Activity name	Activity type	Run start ↑↓	Duration	Status	Integration runtime
Calculate Top 5 Products	SynapseNotebo	11/26/20, 3:39:57 AM	00:05:45	✓ Succeeded	DefaultIntegrationRuntime (East US 2)
write_user_profile_to_asa	ExecuteDataFlov	11/26/20, 3:33:50 AM	00:06:07	✓ Succeeded	DefaultIntegrationRuntime (East US 2)

The Pipeline run ID is highlighted.

These values match because we passed in the Pipeline run ID to the *runId* parameter on the Notebook activity.