Exercise: Join and transform shared and current dataset with Mapping Flow in Azure Data Factory

Transform data using mapping data flow

Now that you have successfully copied data into Azure Data Lake Storage, it is time to join and aggregate that data into a data warehouse.

We will use mapping data flow, Azure Data Factory's visually designed transformation service. Mapping data flows allow users to develop transformation logic code-free and execute them on spark clusters managed by the ADF service.

The data flow created in this step inner joins the 'TripDataCSV' dataset created in the previous section with the 'TripFares.csv' file shared via Azure Data Share. based on four key columns. Then the data gets aggregated based upon column payment_type to calculate the average of certain fields and written in a Azure Synapse Analytics table.

Add a data flow activity to your pipeline

In the activities pane of the pipeline canvas, open the Move and Transform accordion and drag the Data flow activity onto the canvas.

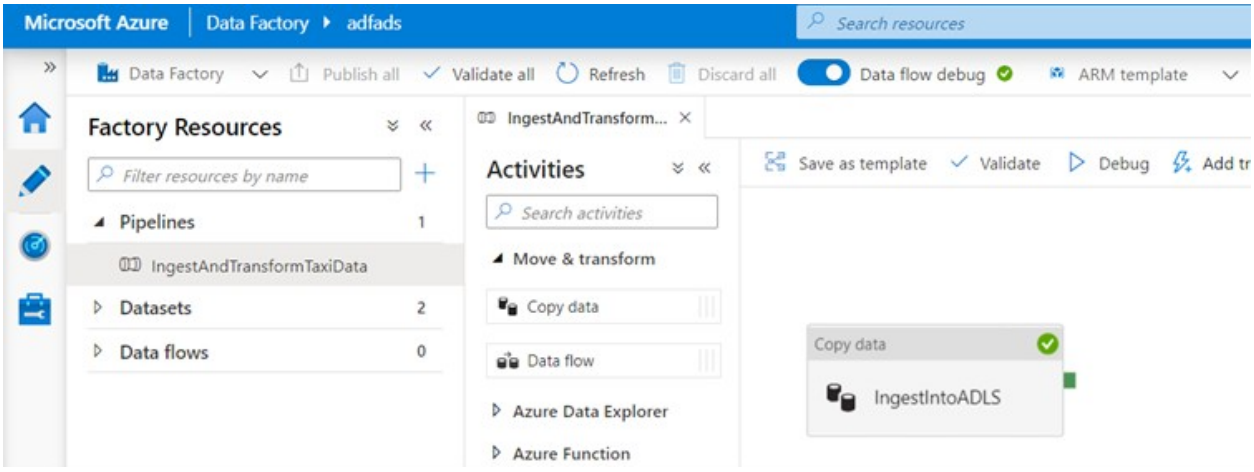

Move and Transform

In the side pane that opens, select Create new data flow and choose Mapping data flow.

Click OK.

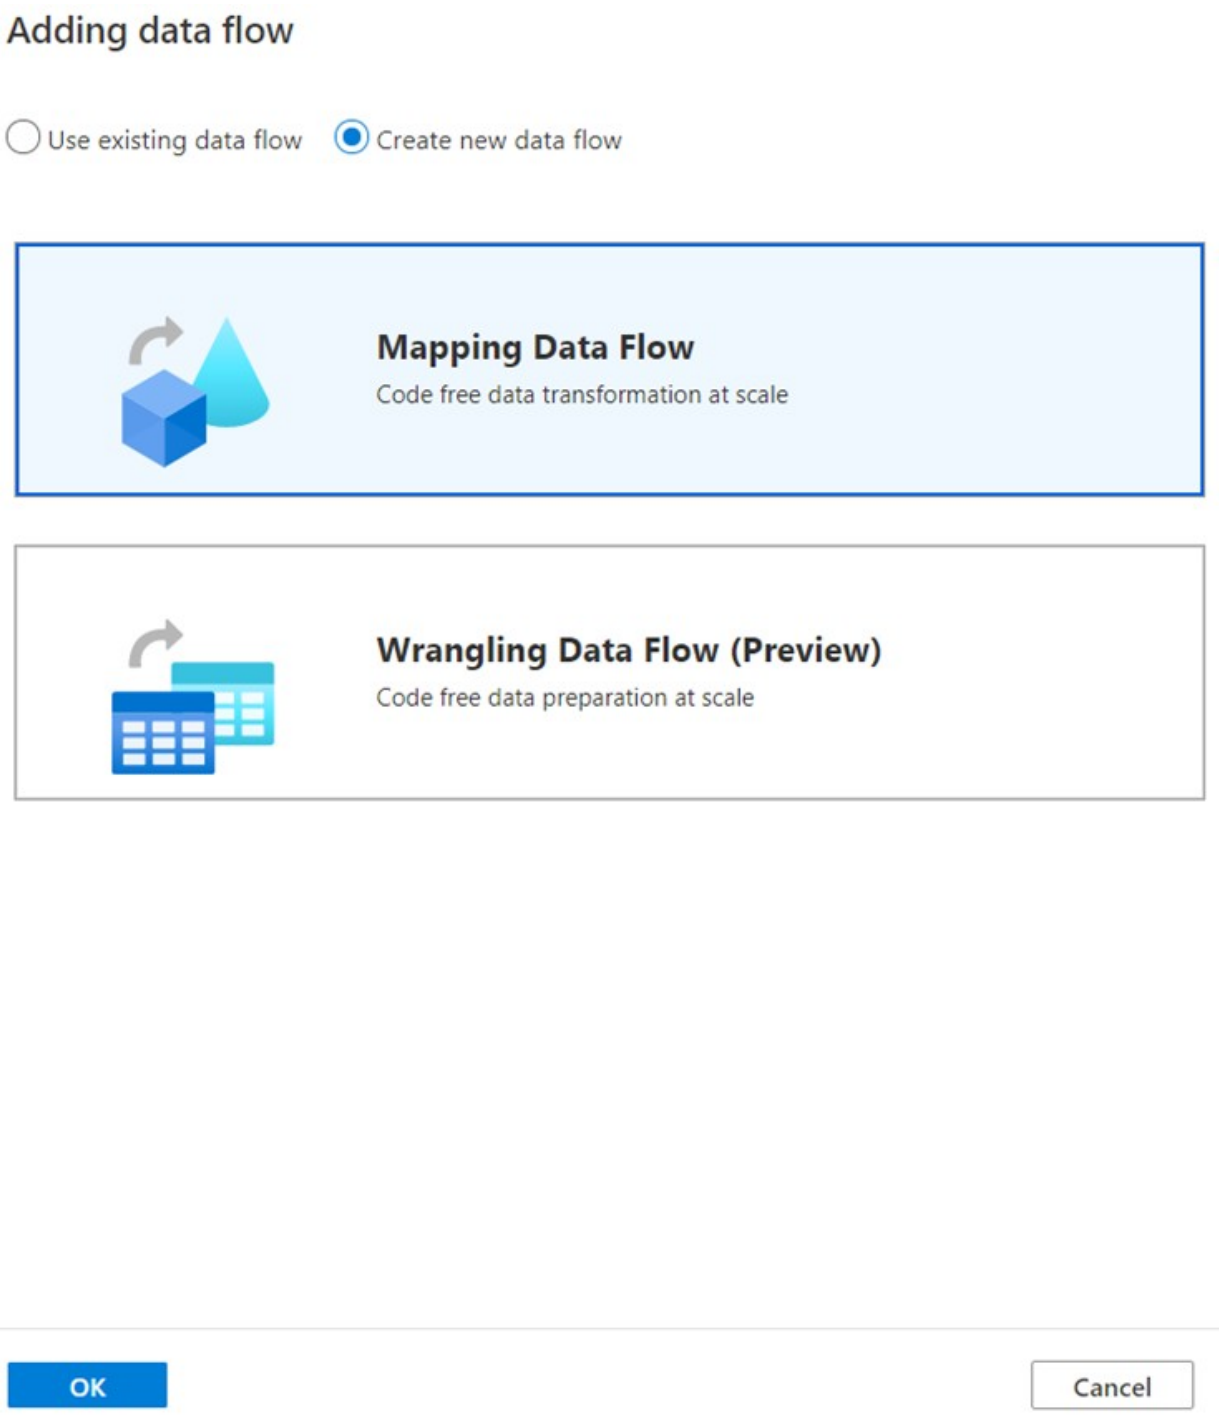

Create new DataFlow

You'll be directed to the data flow canvas where you will be building your transformation logic.

In the general tab, name your data flow 'JoinAndAggregateData'.

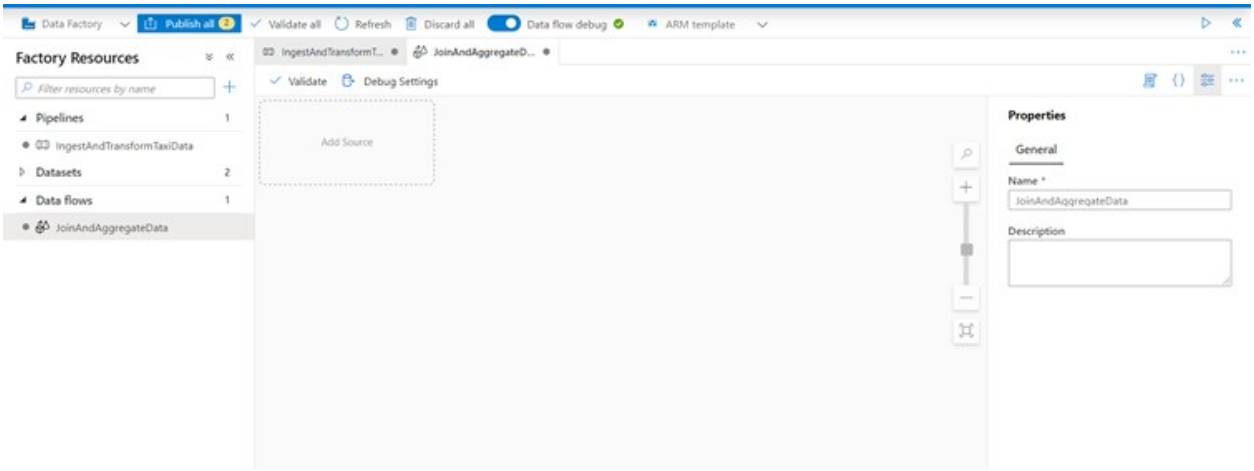

Properties new DataFlow

Configure your trip data csv source

The first thing you want to do is configure your two source transformations. The first source will point to the 'TripDataCSV' DelimitedText dataset. To add a source transformation, click on the Add Source box in the canvas.

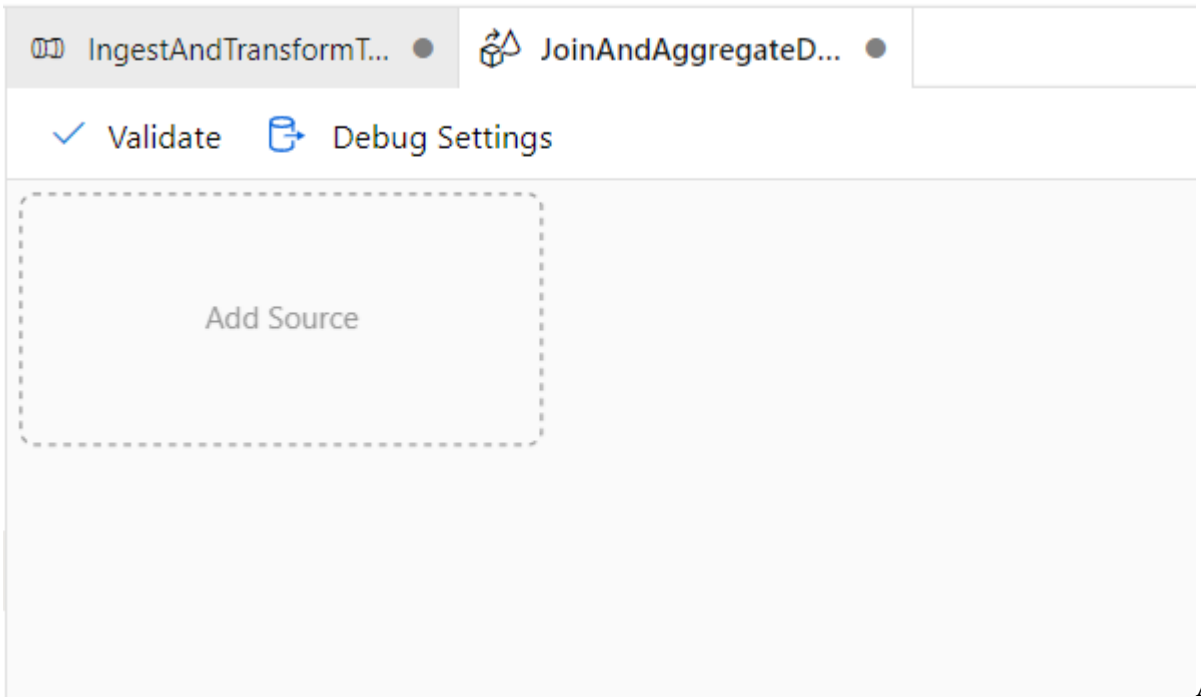

Add Source to Dataflow

Name your source 'TripDataCSV' and select the 'TripDataCSV' dataset from the source drop-down.

Since trip-data.csv exists now, click Open to go to the dataset settings tab.

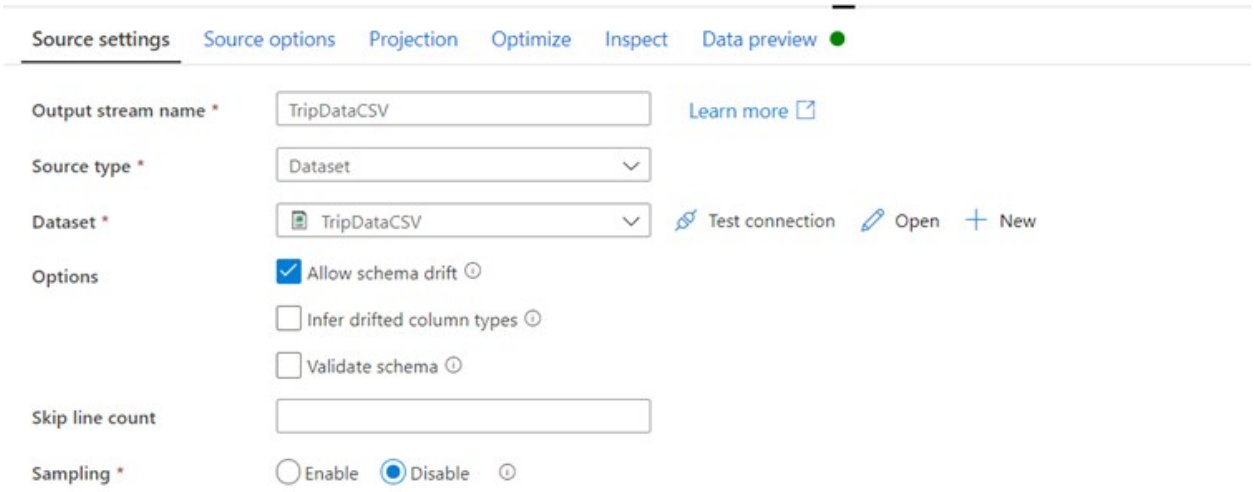

Source Settings of Dataflow

Go to tab Schema and click Import schema.

Select From connection/store to import directly from the file store.

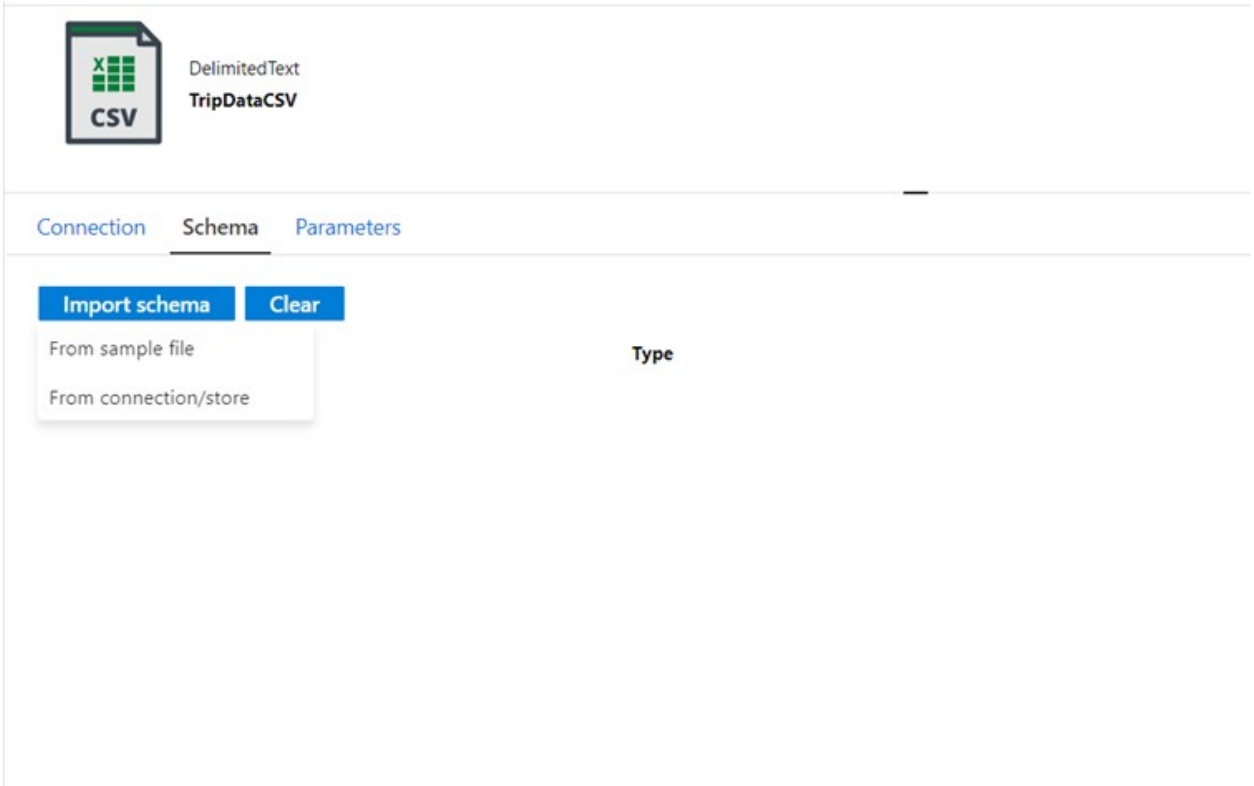

Schema definition source Dataflow

If you select From connection/store the columns should appear as following:

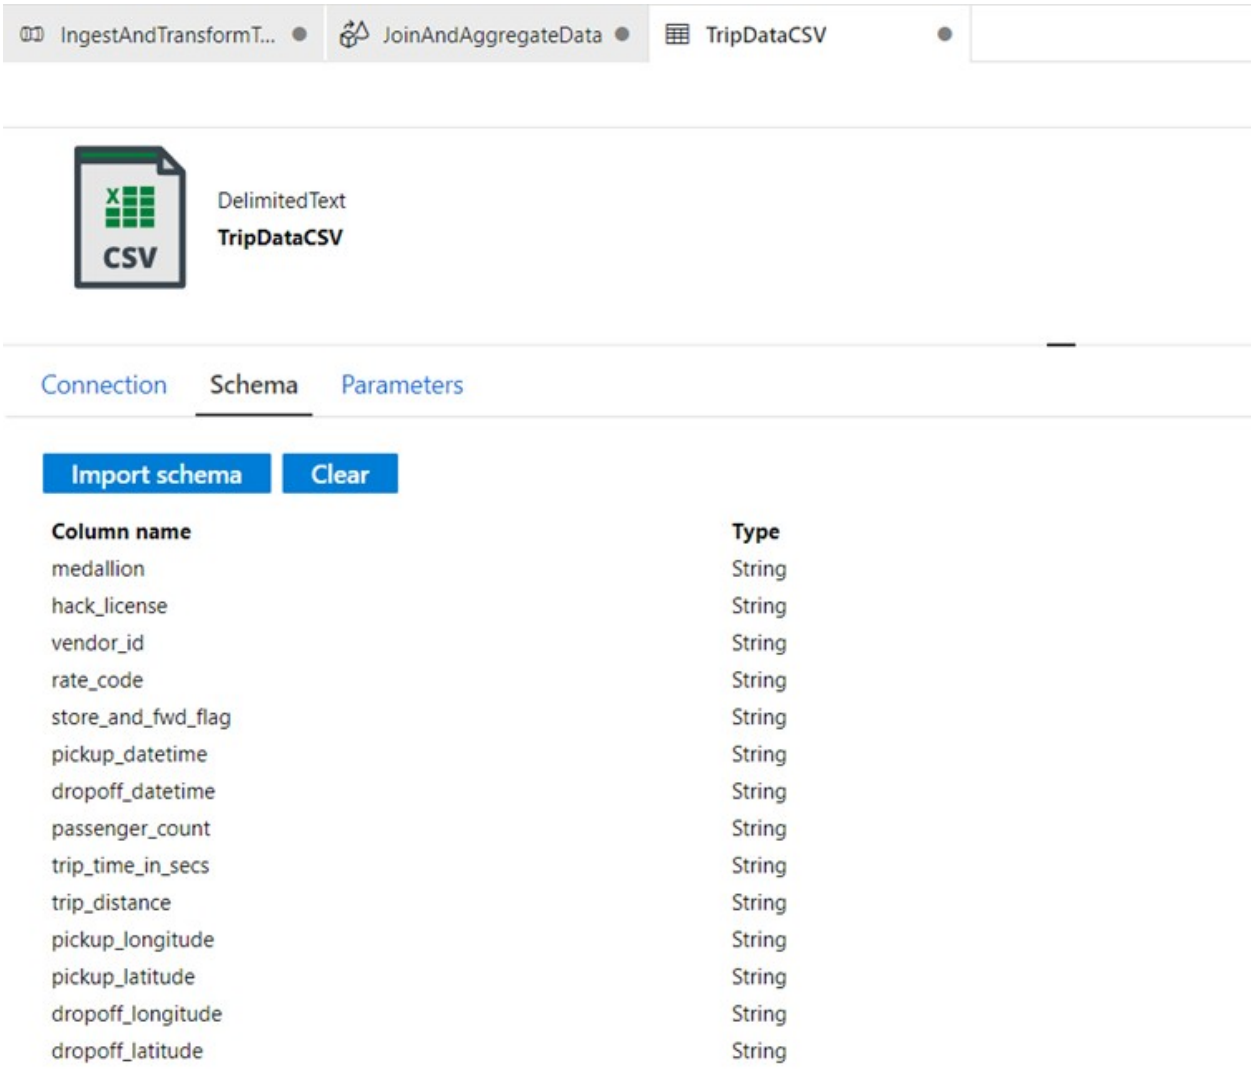

image-07 Schema definition from Connection

Go back to data flow 'JoinAndAggregateData'.

If your debug cluster has started (indicated by a green circle next to the debug slider), you can get a snapshot of the data in the Data Preview tab.

Click Refresh to fetch a data preview.

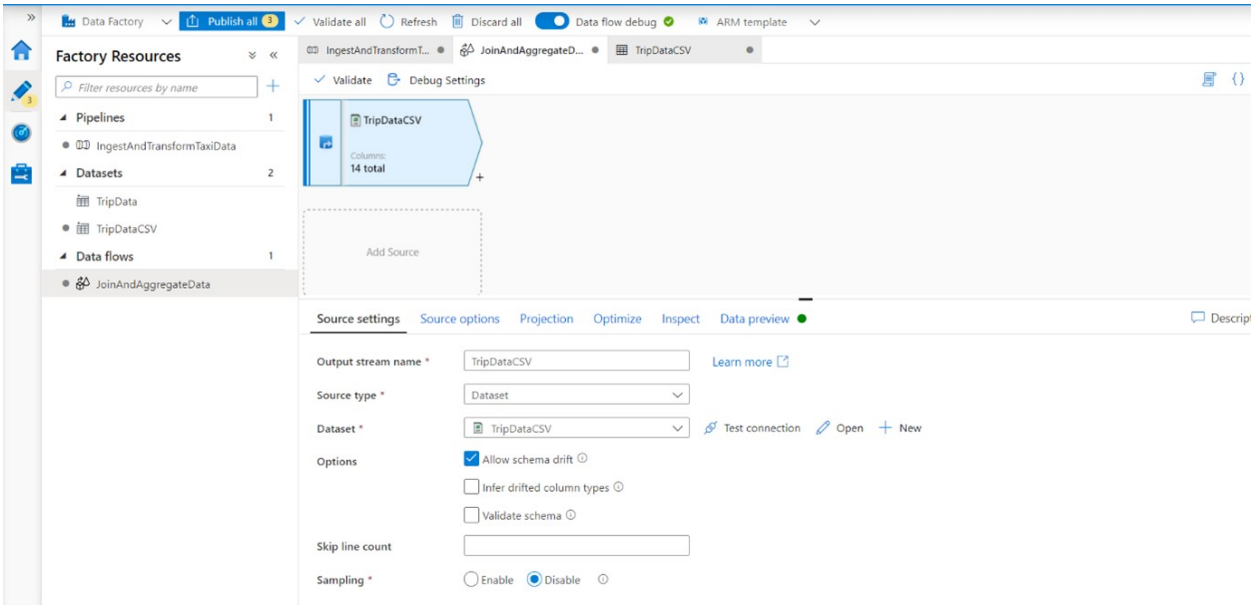

Source Settings after Schema Mapping Dataflow

When you select the Data Preview tab and click refresh the data will be fetched as long as your debug cluster is green, however Data preview does not write data.

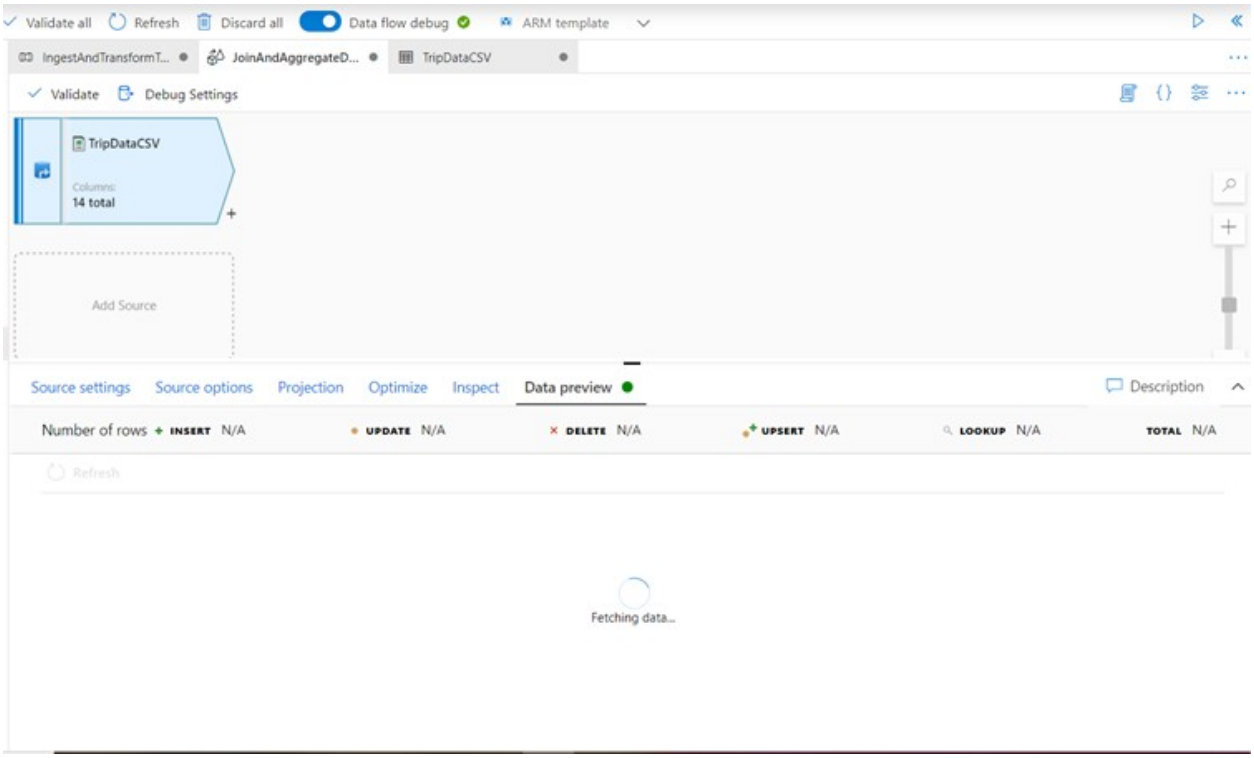

Data Preview Fetching Data

Configure your trip fares csv source:

The second source you're adding will point at the csv file 'TripFares.csv' that was shared during the Azure Data Share exercise.

Under your 'TripDataCSV' source, there will be another Add Source box.

Click it to add a new source transformation.

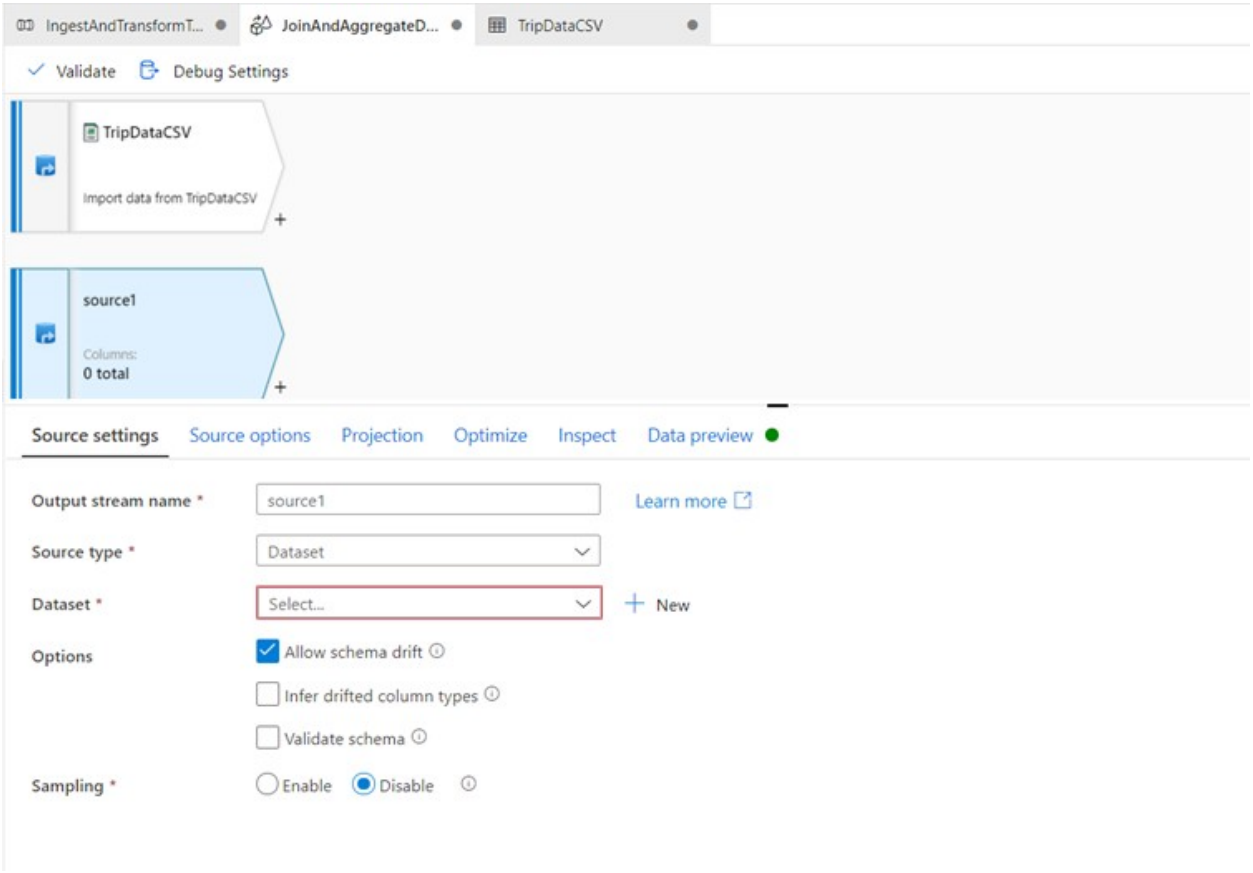

Add Source Transformation

Name this source 'TripFaresCSV'.

Click New next to the source dataset field to create a new ADLS gen2 dataset.

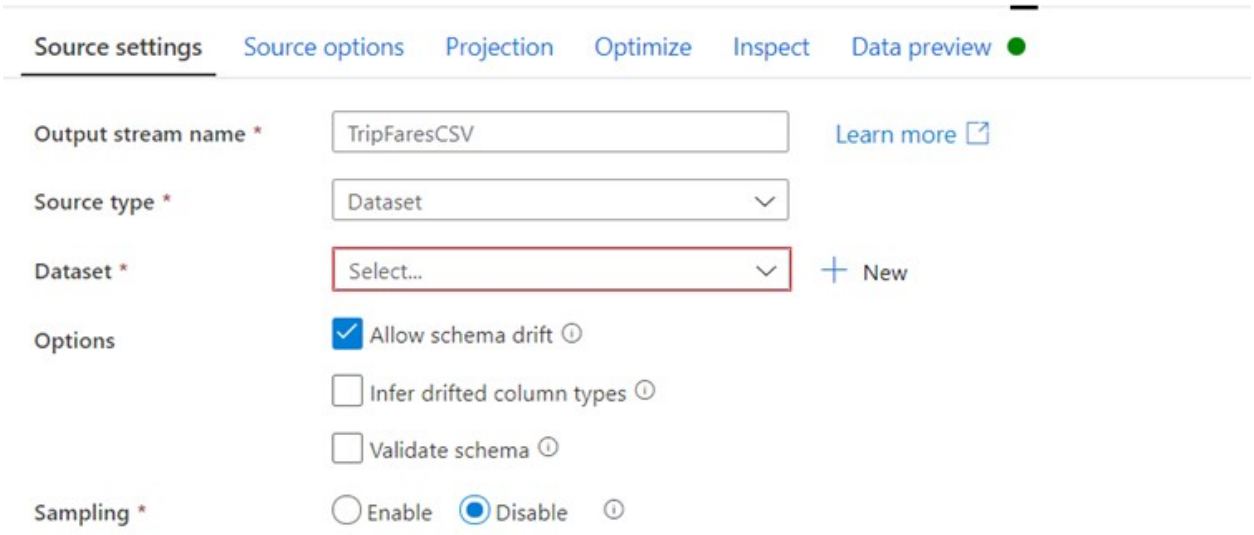

New Source creation for Transformation

Select the Azure Data Lake Storage gen2 tile and click continue.

You may notice many of the connectors in data factory are not supported in mapping data flow. To transform data from one of these sources, ingest it into a supported source using the copy activity.

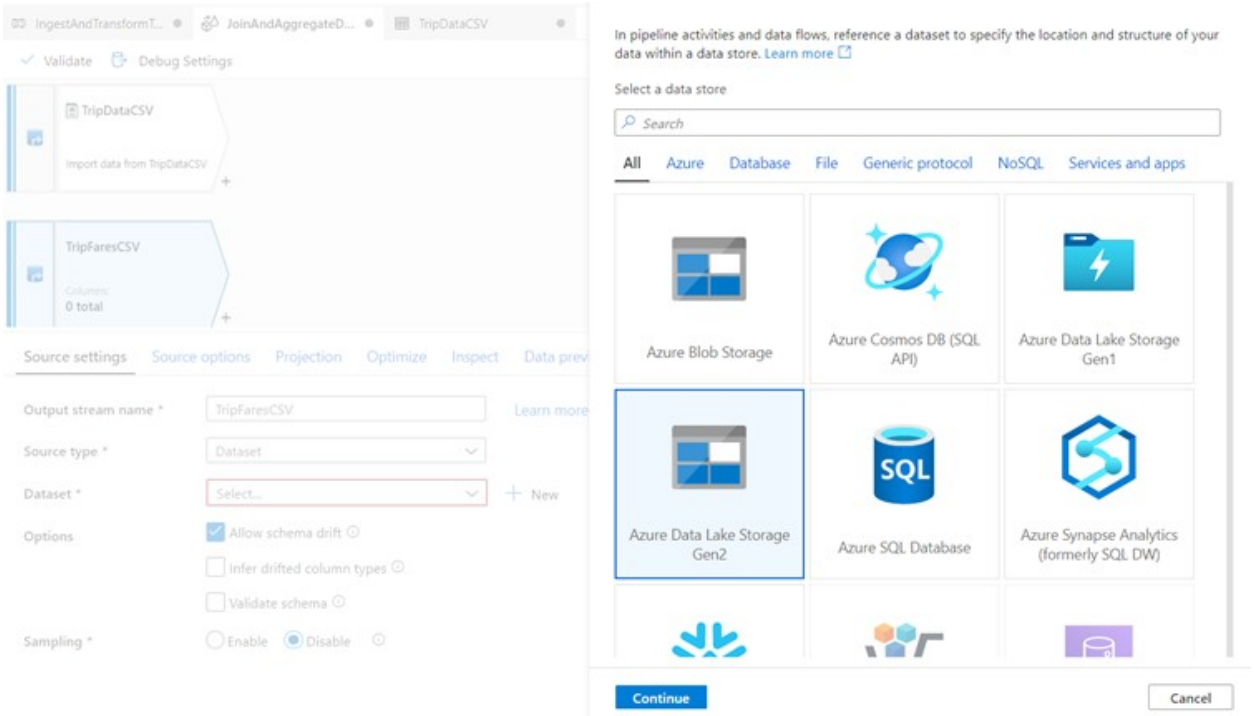

Select Azure Data Lake Storage Gen 2 for new Source Transformation

In the select format pane, select DelimitedText as you are reading from a csv file. Click continue.

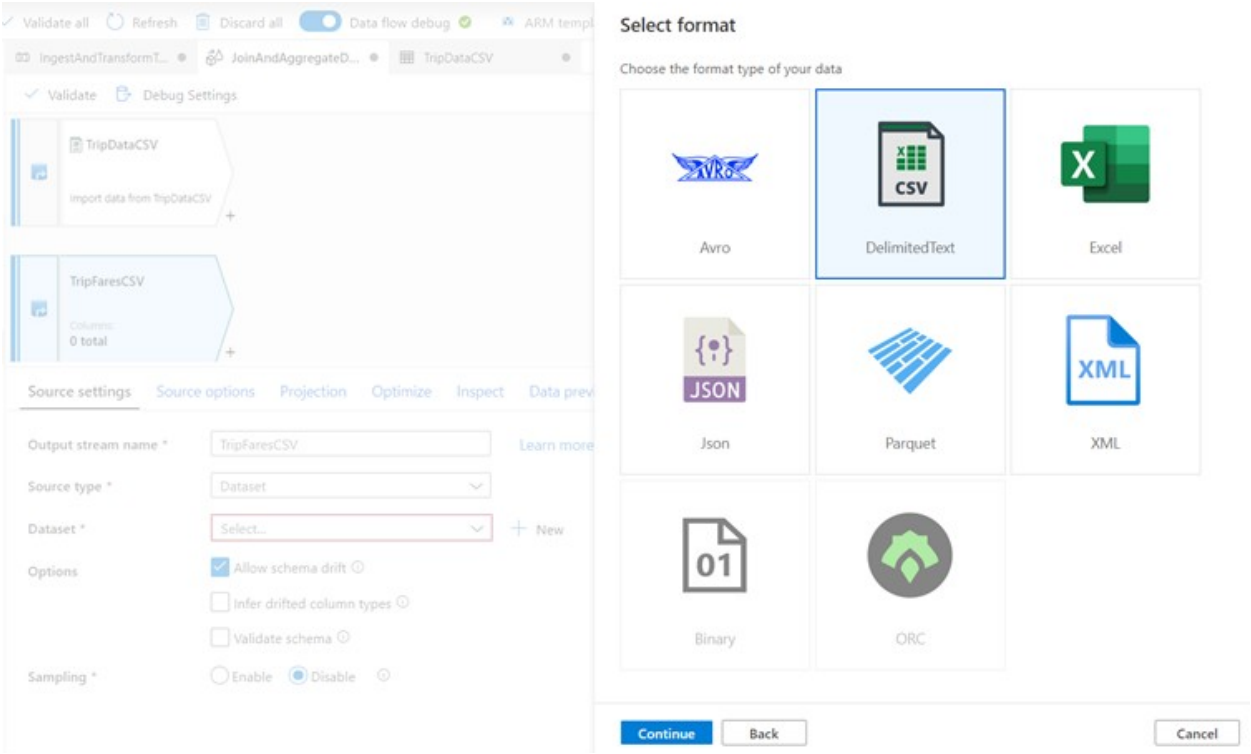

Select Delimited Text

Name your dataset 'TripFaresCSV'.

Select 'ADLSGen2' as your linked service.

Set the file path to taxidata/TripFares.csv. In the Azure Data Share exercise you have called it TripFares.csv in container 'taxidata'.

Set First row as header to true as the input data has headers.

Import the schema From connection/store.

Click OK when finished.

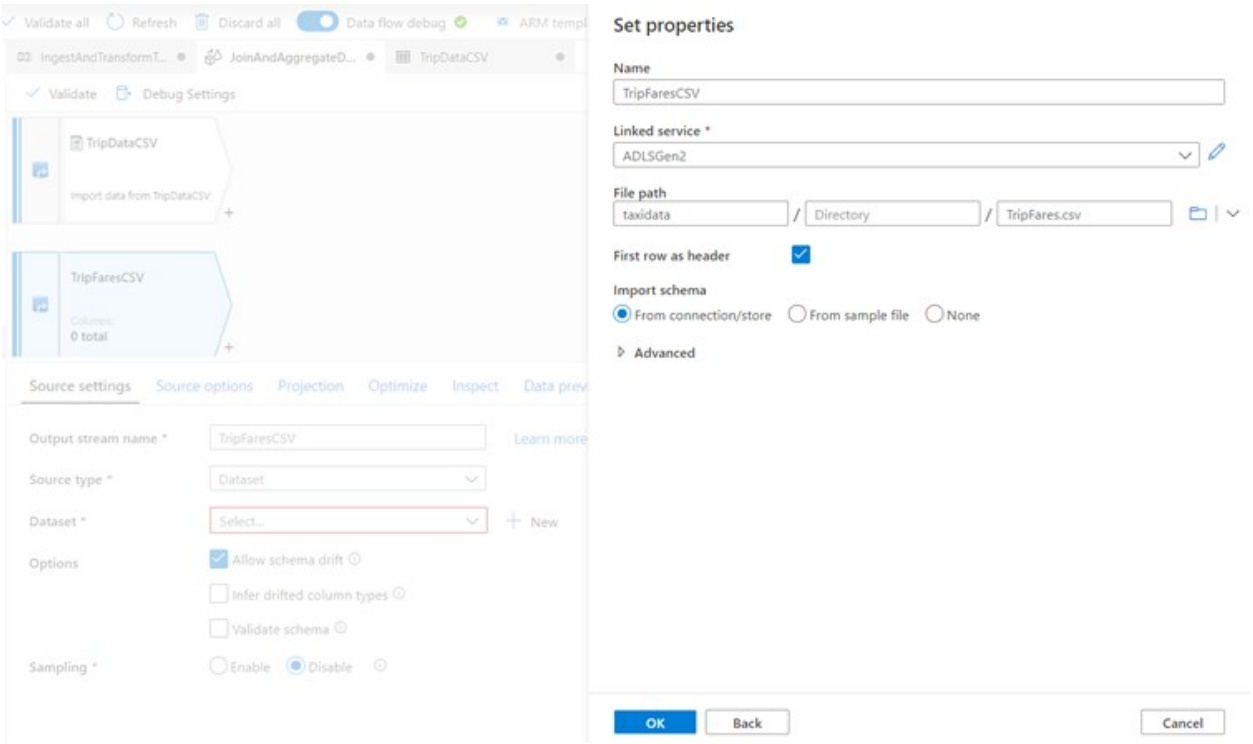

Set properties of File

To verify your source is configured correctly, fetch a data preview in the Data Preview tab.

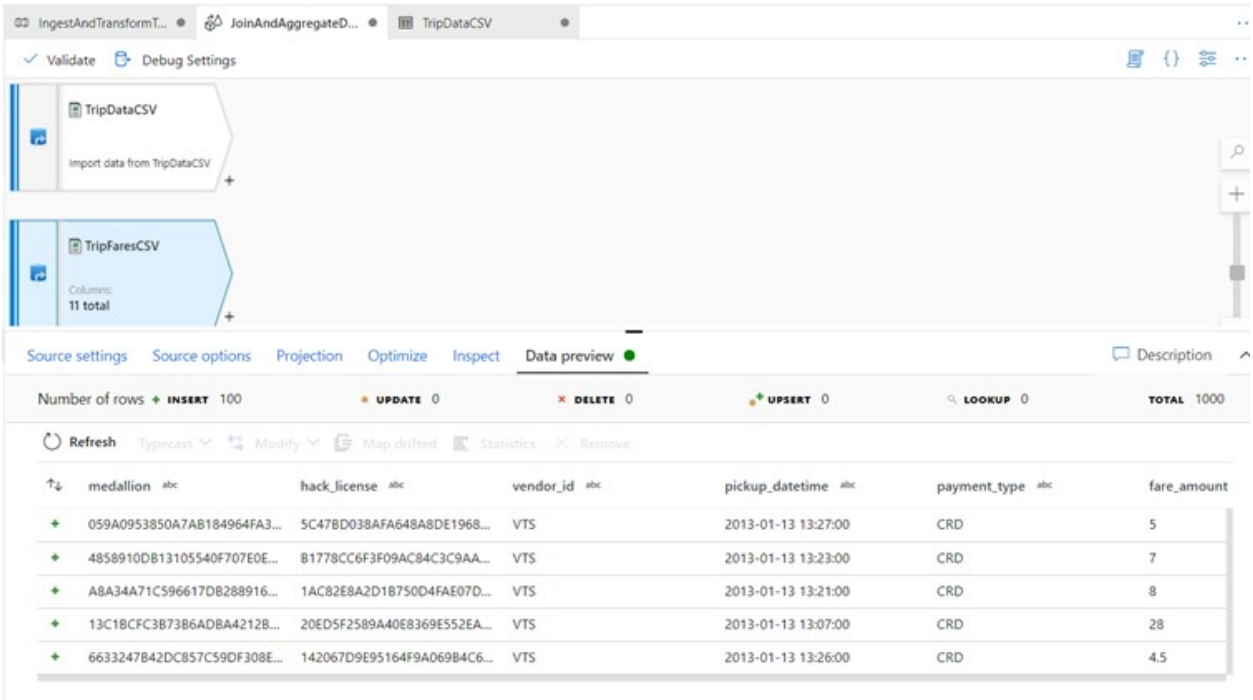

Data Preview of created dataset

Inner join TripDataCSV and TripFaresSQL

To add a new transformation, click the plus icon in the bottom-right corner of 'TripDataCSV'.

Under Multiple inputs/outputs, select Join.

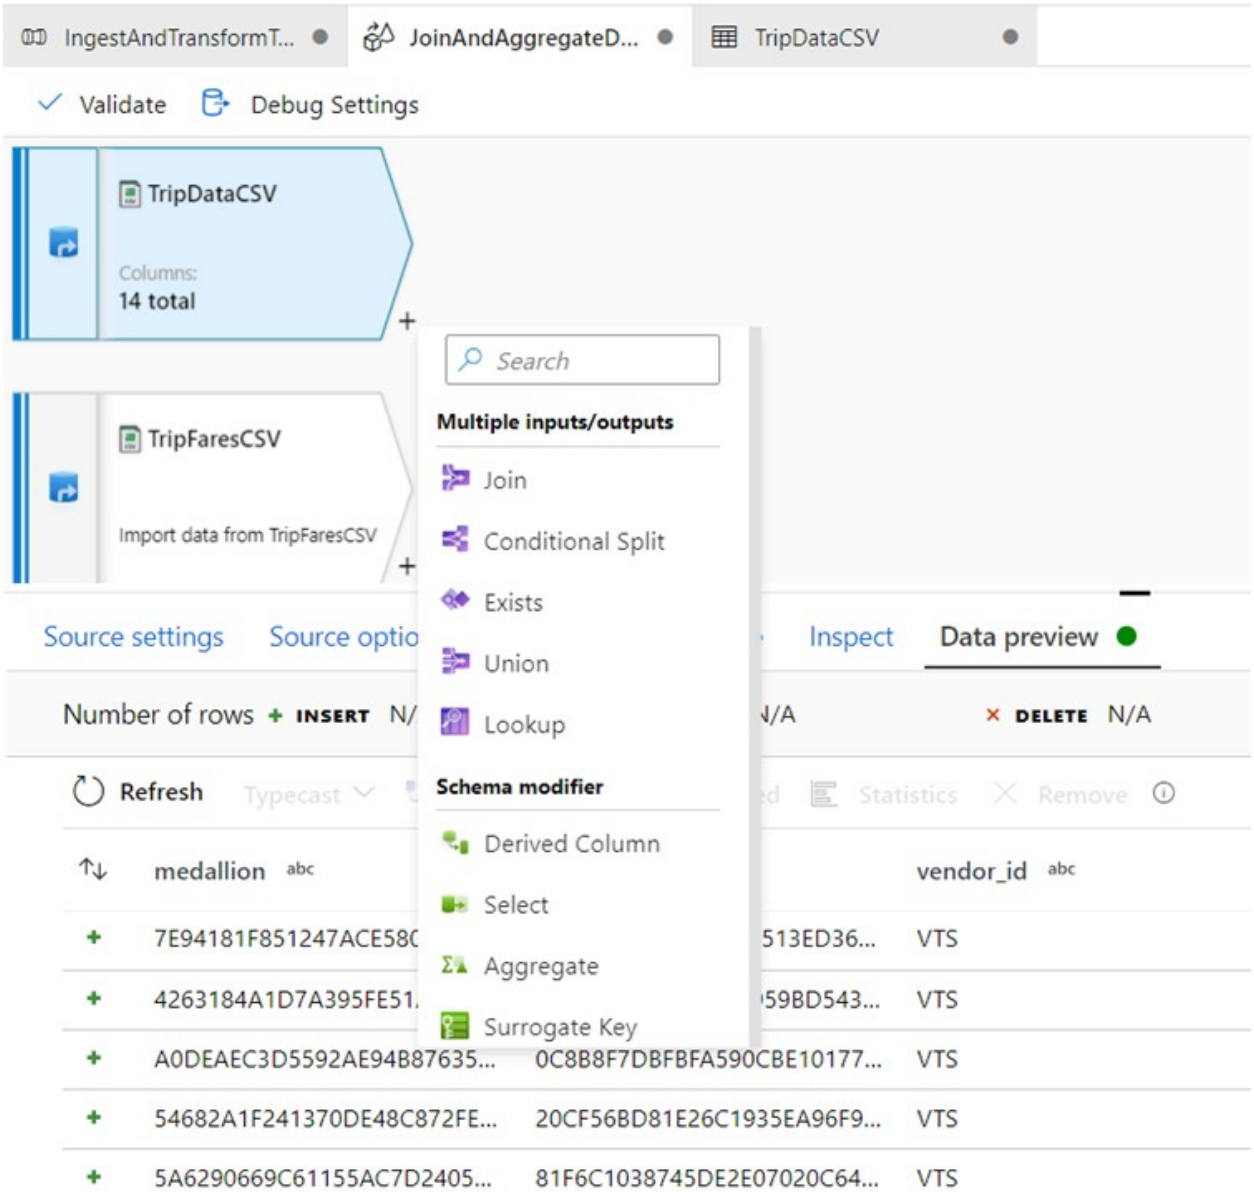

Add Join Transformation

Name your join transformation 'InnerJoinWithTripFares'.

Select 'TripFaresCSV' from the right stream dropdown.

Select Inner as the join type.

Select which columns you wish to match on from each stream via the Join conditions dropdown.

To add an additional join condition, click on the plus icon next to an existing condition.

By default, all join conditions are combined with an AND operator which means all conditions must be met for a match.

We want to match on columns medallion, hack_license, vendor_id, and pickup_datetime

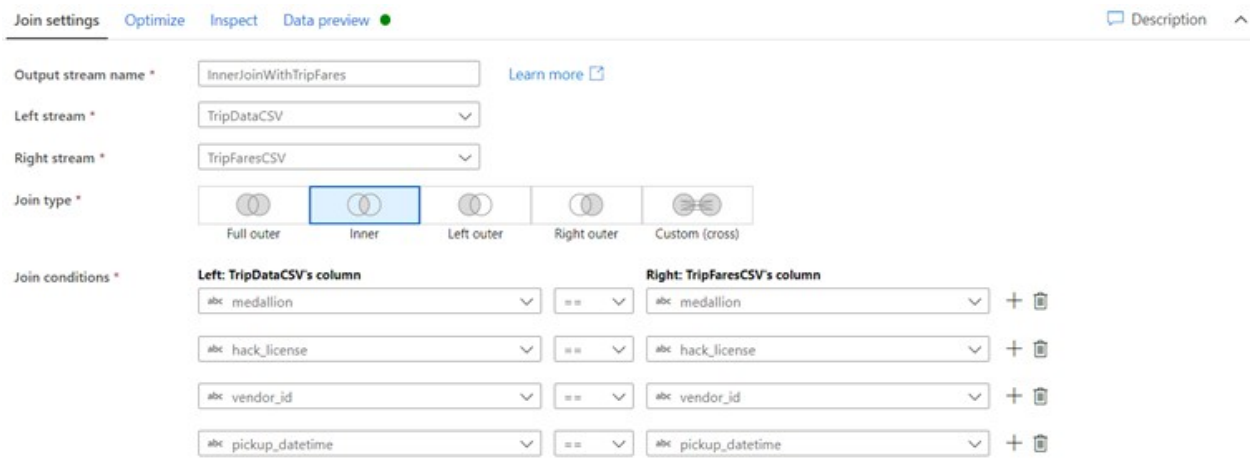

Match Columns for Join Transformation

Verify you successfully joined 25 columns together with Data Preview.

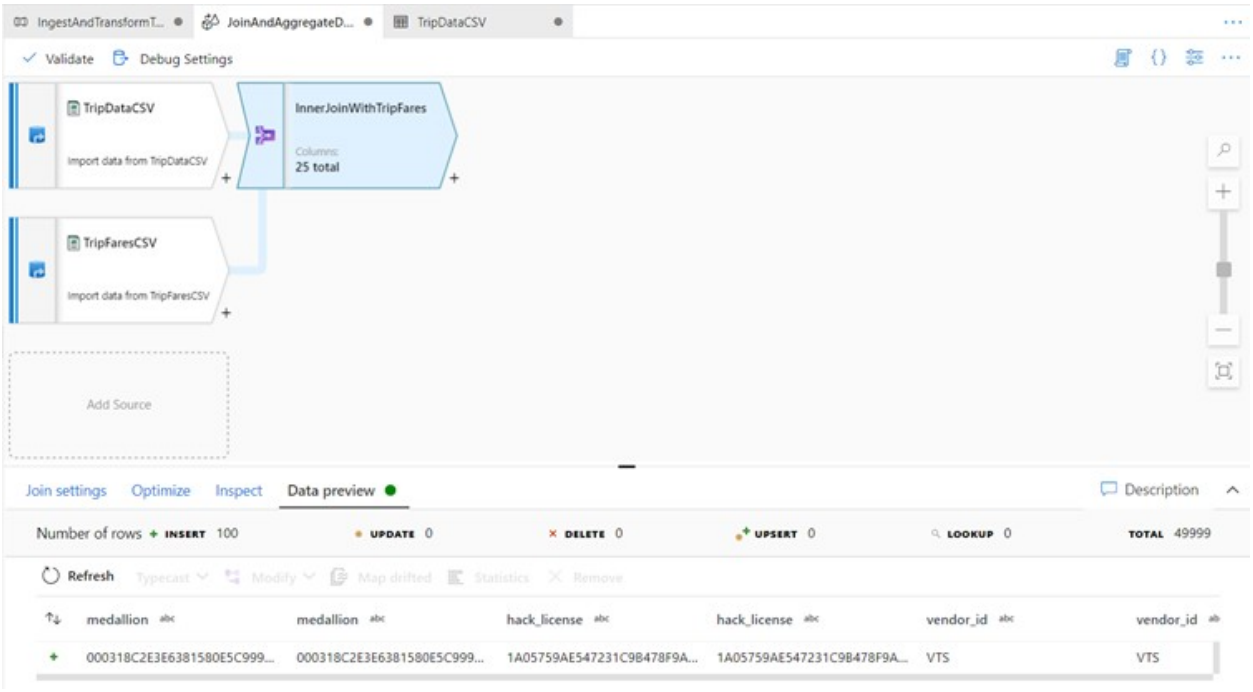

Data Preview for Join Transformation

Aggregate by payment_type

After you complete a join transformation, add an aggregate transformation by clicking the plus icon next to 'InnerJoinWithTripFares.

Choose Aggregate under Schema modifier and call the Aggregate: 'AggregateByPaymentType'

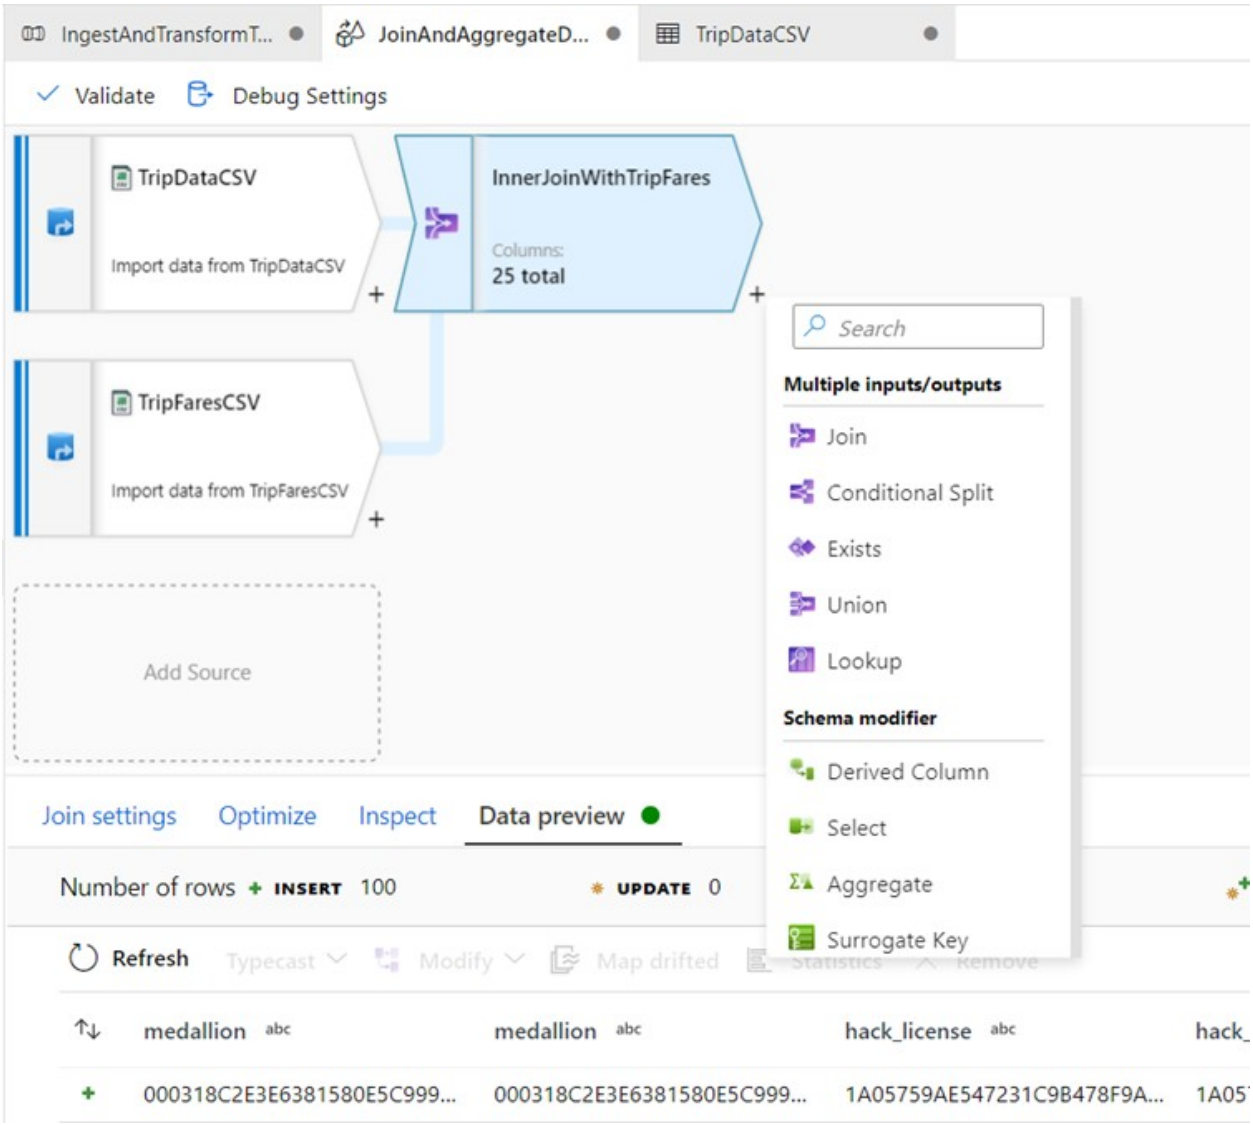

Add Aggregation for Transformation

Name your aggregate transformation 'AggregateByPaymentType'. Select payment_type as the group by column.

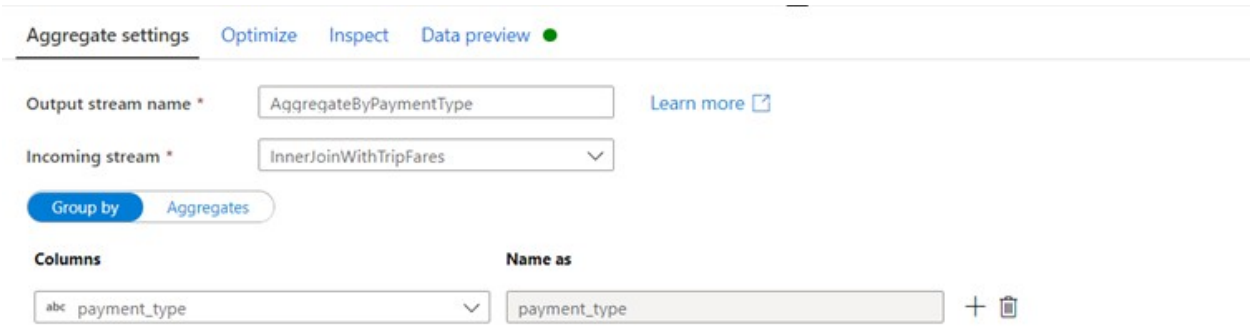

Group by setting for Aggregation

Go to the Aggregates tab. Here, you will specify two aggregations:

```
* The average fare grouped by payment type

* The total trip distance grouped by payment type
```

Create the average fare expression.

In the text box labeled Add or select a column, enter 'average_fare'.

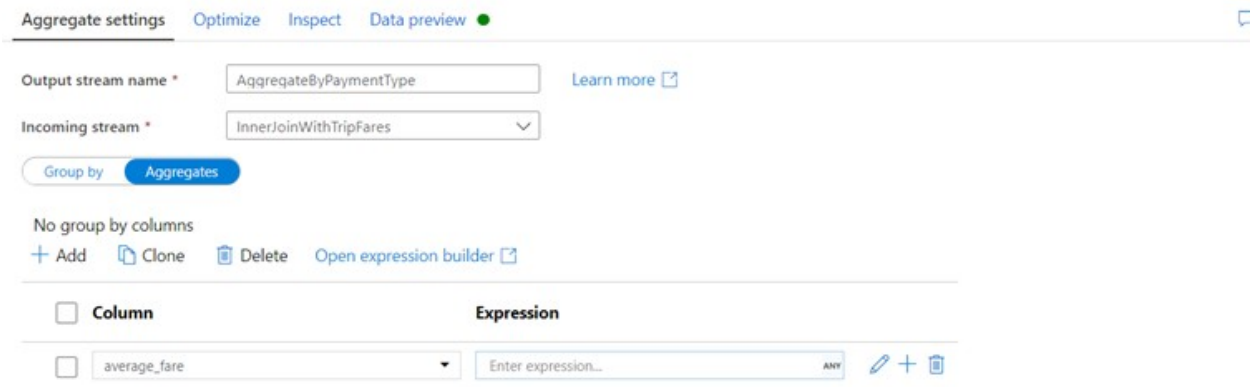

Aggregates settings

To enter an aggregation expression, click on the marker next to the the blue box labeled Enter expression.

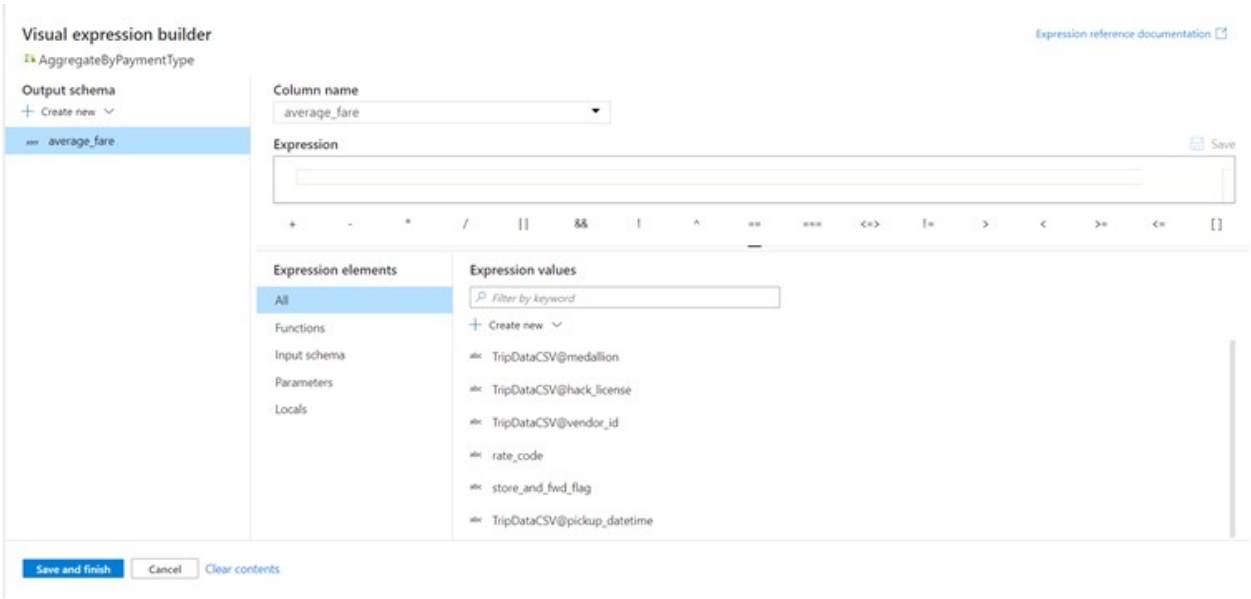

Expression builder for Aggregations

This will open up the data flow expression builder, a tool used to visually create data flow expressions using input schema, built-in functions and operations, and user-defined parameters. For more information on the capabilities of the expression builder, see the expression builder documentation.

To get the average fare, use the avg() aggregation function to aggregate the total_amount column cast to an integer with toInteger(). In the data flow expression language, this is defined as avg(toInteger(total_amount)).

Click Save and finish when you are done.

Visual Expression builder for Aggregations

To add an additional aggregation expression, click on the plus icon next to average_fare.

Select Add column.

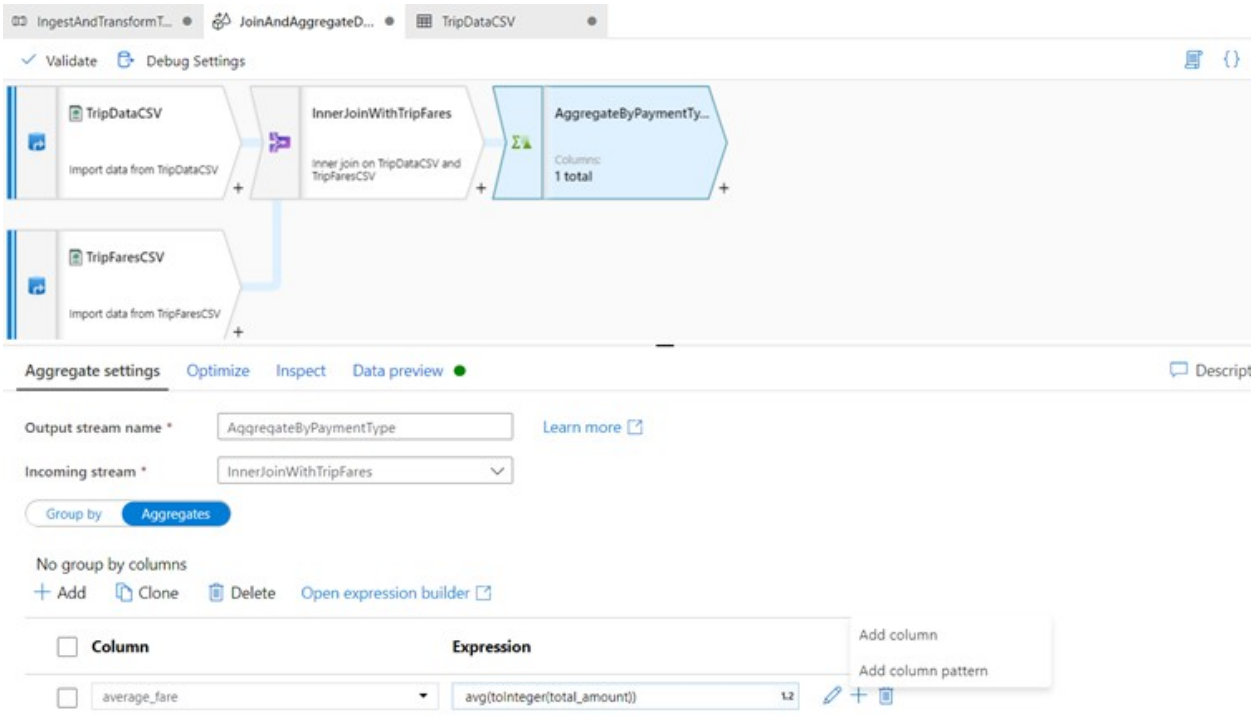

Add a column for another expression

In the text box labeled Add or select a column, enter 'total_trip_distance'. As in the last step, open the expression builder to enter in the expression.

To get the total trip distance, use the sum() aggregation function to aggregate the trip_distance column cast to an integer with toInteger().

In the data flow expression language, this is defined as sum(toInteger(trip_distance)).

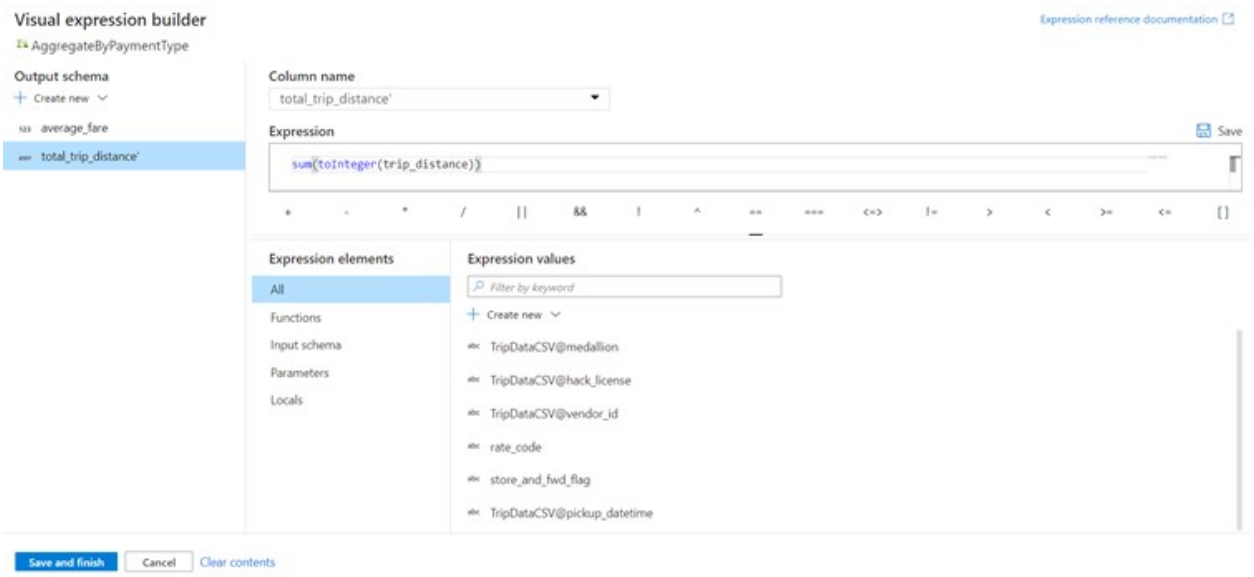

Add expression to new column

Click Save and finish.

Test your transformation logic in the Data Preview tab.

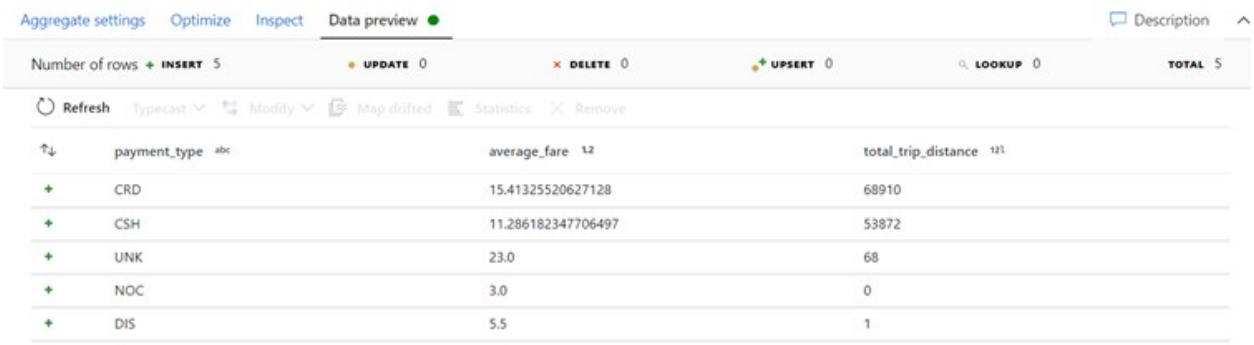

payment_type	average_fare	total_trip_distance
CRD	15.41325520627128	68910
CSH	11.286182347706497	53872
UNK	23.0	68
NOC	3.0	0
DIS	5.5	1